

ECE C247 HW 1 Solution

Ashish Kumar Singh (UID:105479019)

January 16, 2022

Problem 1. Linear Algebra Refresher

Sub-Problem 1(a)(i). Example of A for $AA^T = I$?

Solution 1(a)(i). Let

$$A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$
$$AA^T = \begin{bmatrix} a^2 + b^2 & ac + bd \\ ac + bd & c^2 + d^2 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

To satisfy diagonal terms, we can take $a = \sin \alpha$, $b = \cos \alpha$, $c = \sin \beta$, $d = \cos \beta$
Non-diagonal term becomes,

$$ac + bd = \sin \alpha \sin \beta + \cos \alpha \cos \beta$$

$$0 = \cos(\alpha - \beta)$$

$$\alpha = \beta \pm \frac{(2n+1)\pi}{2}$$

As an example we can take, $\alpha = 30^\circ$ and $\beta = 120^\circ$, we get,

$$A = \begin{bmatrix} \frac{1}{2} & \frac{\sqrt{3}}{2} \\ \frac{\sqrt{3}}{2} & \frac{-1}{2} \end{bmatrix}$$

For eigenvalues and eigenvectors, we set $(A - \lambda I)x = 0$,

$$(A - \lambda I) = \begin{bmatrix} \frac{1}{2} - \lambda & \frac{\sqrt{3}}{2} \\ \frac{\sqrt{3}}{2} & \frac{-1}{2} - \lambda \end{bmatrix}$$

Setting determinant to zero,

$$\frac{(2\lambda-1)(2\lambda+1)}{4} - \frac{3}{4} = 0$$

$$4\lambda^2 = 4$$

$$\lambda = \pm 1$$

For $\lambda = 1$, we can find its corresponding eigenvalue by $(A - I)x = 0$,

$$(A - I)x = \begin{bmatrix} \frac{1}{2} - 1 & \frac{\sqrt{3}}{2} \\ \frac{\sqrt{3}}{2} & -\frac{1}{2} - 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = 0$$

Solving with additional constraint $x_1^2 + x_2^2 = 1$, we get $x_1 = \frac{\sqrt{3}}{2}$ and $x_2 = \frac{1}{2}$, thus eigenvector corresponding to $\lambda = 1$ is:

$$x = \begin{bmatrix} 0.866 \\ 0.5 \end{bmatrix}$$

Similarly, the eigenvector corresponding to $\lambda = -1$ is:

$$x = \begin{bmatrix} 0.5 \\ -0.866 \end{bmatrix}$$

We notice that eigenvalues are unit norm and eigenvectors are orthogonal.

Sub-Problem 1(a)(ii). show eigenvalues are unit norm

Solution 1(a)(ii). If λ is the eigenvalue of A then $Ax = \lambda x$,

$$(Ax)^T(Ax) = (\lambda x)^T(\lambda x)$$

$$x^T A^T A x = |\lambda|^2 x^T x$$

$$x^T x = |\lambda|^2 x^T x$$

$$|\lambda|^2 = 1$$

Sub-Problem 1(a)(iii). show eigenvectors are orthogonal

Solution 1(a)(iii). If λ_1 and λ_2 are two distinct eigenvalues of A then $Ax_1 = \lambda_1 x_1$, $Ax_2 = \lambda_2 x_2$, and $\lambda_1 \neq \lambda_2$

$$(Ax_1)^T(Ax_2) = (\lambda_1 x_1)^T(\lambda_2 x_2)$$

$$x_1^T A^T A x_2 = \lambda_1 \lambda_2 x_1^T x_2$$

$$x_1^T x_2 (\lambda_1 \lambda_2 - 1) = 0$$

We know $|\lambda_1| = 1$, $|\lambda_2| = 1$ and $\lambda_1 \neq \lambda_2$, so $\lambda_1 \lambda_2 \neq 1$ (assuming real eigenvalues), which leaves us with the only option,

$$x_1^T x_2 = 0$$

which means eigenvectors corresponding to distinct eigenvalues of A are orthogonal.

Sub-Problem 1(a)(iv). transformation Ax

Solution 1(a)(iv). Since the eigenvalues are unit norm, the vector x magnitude remains constant, only the direction changes. In other words, vector x is only rotated, its length remain constant.

Sub-Problem 1(b)(i). Relationship between singular vectors and eigenvectors

Solution 1(b)(i). By Singular vector decomposition, we can write, $A = U\Sigma V^T$, where U is the left singular vectors and V is the right singular vectors of A , $UU^T = I$, $V^T V = I$, Σ is a diagonal matrix,

$$\begin{aligned} AA^T &= U\Sigma V^T (U\Sigma V^T)^T \\ AA^T &= U\Sigma V^T V \Sigma^T U^T \\ AA^T &= U\Sigma \Sigma^T U^T \end{aligned}$$

If we multiply AA^T by U , we get,

$$AA^T U = U\Sigma \Sigma^T U^T U = U\Sigma \Sigma^T$$

We can see that the U is the eigenvectors of AA^T , and $\Sigma \Sigma^T$ have the eigenvalues of AA^T . Similarly, for $A^T A$,

$$\begin{aligned} A^T A &= (U\Sigma V^T)^T U\Sigma V^T \\ A^T A &= V \Sigma^T U^T U \Sigma V^T \\ A^T A &= V \Sigma^T \Sigma V^T \end{aligned}$$

If we multiply $A^T A$ by V , we get,

$$A^T A V = V \Sigma^T \Sigma V^T V = V \Sigma^T \Sigma$$

We can see that the V is the eigenvectors of $A^T A$, and $\Sigma^T \Sigma$ have the eigenvalues of $A^T A$.

Hence, the left singular vectors of A are the eigenvectors of AA^T and the right singular vectors of A are the eigenvectors of $A^T A$.

Sub-Problem 1(b)(ii). Relationship between singular values and eigenvalues

Solution 1(b)(ii). Σ is a diagonal matrix, therefore $\Sigma^T \Sigma = \Sigma \Sigma^T$ = square of the diagonal terms of Σ .

Hence, from previous derivation, eigenvalues of AA^T and $A^T A$ is the square of the singular value of A . In other words, singular values of A is square root of eigenvalues of AA^T and $A^T A$.

Sub-Problem 1(c). True or False

Solution 1(c)(i). False

There can be at most n distinct eigenvalues.

Solution 1(c)(ii). False

If eigenvalues are different then it will not be an eigenvector.

Solution 1(c)(iii). True

for eigenvectors, $x^T A x = x^T \lambda x \geq 0$, $\lambda \geq 0$

Solution 1(c)(iv). True

There can be at most n distinct eigenvalues.

Solution 1(c)(v). True

$$A(x + y) = Ax + Ay = \lambda x + \lambda y = \lambda(x + y)$$

Problem 2. Probability Refresher

Sub-Problem 2(a)(i). Find posterior $P(H50|T)$?

Solution 2(a)(i).

$$P(H50|T) = \frac{P(T|H50)P(H50)}{P(T|H50)P(H50) + P(T|H60)P(H60)}$$

$$P(H50|T) = \frac{0.5 * 0.5}{0.5 * 0.5 + 0.4 * 0.5}$$

$$P(H50|T) = \frac{5}{9}$$

Sub-Problem 2(a)(ii). Find posterior $P(H50|THHH)$?

Solution 2(a)(ii).

$$P(H50|THHH) = \frac{P(THHH|H50)P(H50)}{P(THHH|H50)P(H50) + P(THHH|H60)P(H60)}$$

$$P(H50|THHH) = \frac{0.5 * 0.5 * 0.5 * 0.5 * 0.5}{0.5 * 0.5 * 0.5 * 0.5 * 0.5 + 0.4 * 0.6 * 0.6 * 0.6 * 0.5}$$

$$P(H50|THHH) = \frac{0.5 * 0.5 * 0.5 * 0.5 * 0.5}{0.5 * 0.5 * 0.5 * 0.5 * 0.5 + 0.4 * 0.6 * 0.6 * 0.6 * 0.5}$$

$$P(H50|THHH) = \frac{0.03125}{0.03125 + 0.0432} \approx 0.41974$$

Sub-Problem 2(a)(iii). Find posterior for H50,H55,H60 for given data 9 head out of 10 tosses?

Solution 2(a)(iii). Let D denote the data with 9 head out of 10 tosses,

$$P(H50|D) = \frac{P(D|H50)P(H50)}{P(D|H50)P(H50) + P(D|H55)P(H55) + P(D|H60)P(H60)}$$

$$P(H50|D) = \frac{(1/3) * (0.5)^{10}}{(1/3) * (0.5)^{10} + (1/3) * (0.55)^9 * (0.45) + (1/3) * (0.6)^9 * (0.4)}$$

$$P(H50|D) \approx 0.13793$$

$$P(H55|D) = \frac{P(D|H55)P(H55)}{P(D|H50)P(H50) + P(D|H55)P(H55) + P(D|H60)P(H60)}$$

$$P(H55|D) = \frac{(1/3) * (0.55)^9 * (0.45)}{(1/3) * (0.5)^{10} + (1/3) * (0.55)^9 * (0.45) + (1/3) * (0.6)^9 * (0.4)}$$

$$P(H55|D) \approx 0.29271$$

$$P(H60|D) = \frac{P(D|H60)P(H60)}{P(D|H50)P(H50) + P(D|H55)P(H55) + P(D|H60)P(H60)}$$

$$P(H60|D) = \frac{(1/3) * (0.6)^9 * (0.4)}{(1/3) * (0.5)^{10} + (1/3) * (0.55)^9 * (0.45) + (1/3) * (0.6)^9 * (0.4)}$$

$$P(H60|D) \approx 0.56936$$

Sub-Problem 2(b). Find $P(\text{pregnant}|\text{positive})$?

Solution 2(b). We can infer following probabilities:

$P(\text{positive}|\text{pregnant}) = 0.99$, $P(\text{positive}|\text{notpregnant}) = 0.1$,

$P(\text{pregnant}) = 0.01$, $P(\text{notpregnant}) = 0.1$

$$P(\text{pregnant}|\text{positive}) = \frac{P(\text{positive}, \text{pregnant})}{P(\text{positive}, \text{pregnant}) + P(\text{positive}, \text{notpregnant})}$$

$$P(\text{pregnant}|\text{positive}) = \frac{0.99 * 0.01}{0.99 * 0.01 + 0.99 * 0.1}$$

$$P(\text{pregnant}|\text{positive}) = \frac{1}{11} \approx 0.091$$

The answer shows that the test is very bad. It makes sense because the test is giving 10% false positive on a 99% of non-pregnant female population. So, e.g. for a female population of 1000, 990 are not pregnant and only 10 are pregnant but test gives 99 female as false positive. Thus, very low $P(\text{pregnant}|\text{positive})$ makes sense, the test is very bad.

Sub-Problem 2(c). Find $\mathbb{E}(Ax + b)$?

Solution 2(c). Lets find the i^{th} term of $\mathbb{E}(Ax)$,

$$\mathbb{E}(Ax)_i = \mathbb{E}(\sum A_{i,j}x_j)$$

$$\mathbb{E}(Ax)_i = \sum A_{i,j} \mathbb{E}(x)_j$$

$$\mathbb{E}(Ax)_i = (A \mathbb{E}(x))_i$$

The above is true for every i , so we get,

$$\mathbb{E}(Ax + b) = \mathbb{E}(Ax) + b$$

$$\mathbb{E}(Ax + b) = A \mathbb{E}(x) + b$$

Sub-Problem 2(d). Find $cov(Ax + b)$?

Solution 2(d).

$$\text{cov}(Ax + b) = \mathbb{E} \left(((Ax + b) - \mathbb{E}(Ax + b))((Ax + b) - \mathbb{E}(Ax + b))^T \right)$$

$$\text{cov}(Ax + b) = \mathbb{E} \left((A(x - \mathbb{E}(x)))(A(x - \mathbb{E}(x)))^T \right)$$

$$\text{cov}(Ax + b) = \mathbb{E} \left(A(x - \mathbb{E}(x))(x - \mathbb{E}(x))^T A^T \right)$$

$$\text{cov}(Ax + b) = A \mathbb{E} \left((x - \mathbb{E}(x))(x - \mathbb{E}(x))^T \right) A^T$$

$$\text{cov}(Ax + b) = A \text{cov}(x) A^T$$

Problem 3. Multivariate derivatives

Solution 3(a). First finding derivative wrt to each x_i , then combining

$$\nabla_{x_i} x^T A y = \nabla_{x_i} \left(\sum_{i=1}^n \sum_{j=1}^m x_i A_{ij} y_j \right)$$

$$\nabla_{x_i} x^T A y = \sum_{j=1}^m A_{ij} y_j$$

$$\nabla_x x^T A y = \begin{bmatrix} \sum_{j=1}^m A_{1j} y_j \\ \vdots \\ \sum_{j=1}^m A_{nj} y_j \end{bmatrix}$$

$$\nabla_x x^T A y = A y$$

Solution 3(b). First finding derivative wrt to each y_j , then combining

$$\nabla_{y_j} x^T A y = \nabla_{y_j} \left(\sum_{i=1}^n \sum_{j=1}^m x_i A_{ij} y_j \right)$$

$$\nabla_{y_j} x^T A y = \sum_{i=1}^n x_i A_{ij}$$

$$\nabla_{y_j} x^T A y = \sum_{i=1}^n A_{ji}^T x_i$$

$$\nabla_y x^T A y = \begin{bmatrix} \sum_{i=1}^n A_{1i}^T x_i \\ \vdots \\ \sum_{i=1}^n A_{mi}^T x_i \end{bmatrix}$$

$$\nabla_y x^T A y = A^T x$$

Solution 3(c). First finding derivative wrt to each A_{ij} , then combining

$$\nabla_{A_{ij}} x^T A y = \nabla_{A_{ij}} \left(\sum_{i=1}^n \sum_{j=1}^m x_i A_{ij} y_j \right)$$

$$\nabla_{A_{ij}} x^T A y = x_i y_j$$

$$\nabla_A x^T A y = \begin{bmatrix} x_1 y_1 & \dots & x_1 y_m \\ \vdots & \vdots & \vdots \\ x_n y_1 & \dots & x_n y_m \end{bmatrix}$$

$$\nabla_A x^T A y = x y^T$$

Solution 3(d). First finding derivative wrt to each x_k , then combining

$$\nabla_{x_k} \left(x^T A x + b^T x \right) = \nabla_{x_k} \left(\sum_{i=1}^n \sum_{j=1}^m A_{ij} x_i x_j + \sum_{i=1}^n b_i x_i \right)$$

$$\nabla_{x_k} \left(x^T A x + b^T x \right) = \sum_{j=1}^m A_{kj} x_j + \sum_{i=1}^n A_{ik} x_i + b_i$$

$$\nabla_{x_k} \left(x^T A x + b^T x \right) = \sum_{j=1}^m A_{kj} x_j + \sum_{i=1}^n A_{ki}^T x_i + b_i$$

$$\nabla_x \left(x^T A x + b^T x \right) = \begin{bmatrix} \sum_{j=1}^m A_{1j} x_j + \sum_{i=1}^n A_{1i}^T x_i + b_1 \\ \vdots \\ \sum_{j=1}^m A_{nj} x_j + \sum_{i=1}^n A_{ni}^T x_i + b_n \end{bmatrix}$$

$$\nabla_x \left(x^T A x + b^T x \right) = A x + A^T x + b$$

Solution 3(e). Lets assume, dimension of A and B are $(n \times m)$ and $(m \times n)$ respectively. First finding derivative wrt to each A_{ij} , then combining

$$\nabla_{A_{ij}} \text{tr}(AB) = \nabla_{A_{ij}} \left(\sum_{i=1}^n (AB)_{ii} \right)$$

$$\nabla_{A_{ij}} \text{tr}(AB) = \nabla_{A_{ij}} \left(\sum_{i=1}^n \sum_{j=1}^m A_{ij} B_{ji} \right)$$

$$\nabla_{A_{ij}} \text{tr}(AB) = B_{ji} = B_{ij}^T$$

$$\nabla_{A_{ij}} \text{tr}(AB) = \begin{bmatrix} B_{11}^T & \dots & B_{1m}^T \\ \vdots & \vdots & \vdots \\ B_{n1}^T & \dots & B_{nm}^T \end{bmatrix}$$

$$\nabla_A \text{tr}(AB) = B^T$$

Problem 4. Least squares derivation

Solution 4. We can stack all the $y^{(i)}$ as column to form Y , similarly for $x^{(i)}$ to form X , then loss can be written as:

$$L(W) = \frac{1}{2} \sum_{i=1}^n \|y^{(i)} - Wx^{(i)}\|^2 = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - Wx^{(i)})^T (y^{(i)} - Wx^{(i)})$$

$$L(W) = \frac{1}{2} \text{tr} \left((Y - WX)^T (Y - WX) \right)$$

$$L(W) = \frac{1}{2} \text{tr} \left(Y^T Y - Y^T W X - X^T W^T Y + X^T W^T W X \right)$$

Using trace properties: $\text{tr}(A) = \text{tr}(A^T)$, $\text{tr}(A + B) = \text{tr}(A) + \text{tr}(B)$ and $\text{tr}(ABC) = \text{tr}(BCA) = \text{tr}(CAB)$, we get,

$$L(W) = \frac{1}{2} \left(\text{tr}(Y^T Y) - \text{tr}(Y^T W X) - \text{tr}(X^T W^T Y) + \text{tr}(X^T W^T W X) \right)$$

$$L(W) = \frac{1}{2} \left(\text{tr}(Y^T Y) - \text{tr}(W X Y^T) - \text{tr}(W X Y^T) + \text{tr}(W X X^T W^T) \right)$$

Differentiating wrt to W , and setting it to zero to get W_{opt} ,

$$\nabla_W L(W) = \frac{1}{2} \left(0 - Y X^T - Y X^T + W X X^T + W X X^T \right) = 0$$

$$-Y X^T + W_{opt} X X^T = 0$$

$$W_{opt} X X^T = Y X^T$$

$$W_{opt} = Y X^T (X X^T)^{-1}$$

Problem 5. Hello World in Jupyter

Solution 5. code is attached below

Linear regression workbook

This workbook will walk you through a linear regression example. It will provide familiarity with Jupyter Notebook and Python. Please print (to pdf) a completed version of this workbook for submission with HW #1.

ECE C147/C247 Winter Quarter 2022, Prof. J.C. Kao, TAs Y. Li, P. Lu, T. Monsoor, T. wang

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

#allows matlab plots to be generated in line
%matplotlib inline
```


Data generation

For any example, we first have to generate some appropriate data to use. The following cell generates data according to the model: $y = x - 2x^2 + x^3 + \epsilon$

```
In [2]: np.random.seed(0) # Sets the random seed.
num_train = 200 # Number of training data points

# Generate the training data
x = np.random.uniform(low=0, high=1, size=(num_train,))
y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
```

Out[2]: Text(0, 0.5, '\$y\$')



QUESTIONS:

Write your answers in the markdown cell below this one:

- (1) What is the generating distribution of x ?
- (2) What is the distribution of the additive noise ϵ ?

ANSWERS:

- (1) x is generated from a **uniform** distribution in range 0 to 1.
- (2) ϵ noise is generated from a **gaussian** distribution with mean 0 and standard deviation 0.03.

Fitting data to the model (5 points)

Here, we'll do linear regression to fit the parameters of a model $y = ax + b$.

```
In [4]: # xhat = (x, 1)
xhat = np.vstack((x, np.ones_like(x)))

# ===== #
# START YOUR CODE HERE #
# ===== #
# GOAL: create a variable theta; theta is a numpy array whose elements are [a, b]


theta = np.linalg.inv(xhat.dot(xhat.T)).dot(xhat.dot(y))
#print(theta)

# ===== #
# END YOUR CODE HERE #
# ===== #
```

```
In [5]: # Plot the data and your model fit.
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')

# Plot the regression line
xs = np.linspace(min(x), max(x), 50)
xs = np.vstack((xs, np.ones_like(xs)))
plt.plot(xs[0,:], theta.dot(xs))
```

Out[5]: [



QUESTIONS

- (1) Does the linear model under- or overfit the data?
- (2) How to change the model to improve the fitting?

ANSWERS

- (1) The linear model underfit the data because there is non-linearity in data which the linear model cannot fit.
- (2) Instead of linear model we can fit higher degree polynomial models.

Fitting data to the model (10 points)

Here, we'll now do regression to polynomial models of orders 1 to 5. Note, the order 1 model is the linear model you prior fit.

In [6]:

```
N = 5
xhats = []
thetas = []

# ===== #
# START YOUR CODE HERE #
# ===== #

# GOAL: create a variable thetas.
# thetas is a list, where theta[i] are the model parameters for the polynomial fit of order i+1
# i.e., thetas[0] is equivalent to theta above.
# i.e., thetas[1] should be a length 3 np.array with the coefficients of the x^2, x, and 1 re
# ... etc.

for i in range(N):
    if i==0:
        xhats.append(xhat)
    else:
        xhats.append(np.vstack((x**(i+1),xhats[i-1])))
        thetas.append(np.linalg.inv(xhats[i].dot(xhats[i].T)).dot(xhats[i].dot(y)))
#print(thetas)
# ===== #
# END YOUR CODE HERE #
# ===== #
```


In [7]:

```
# Plot the data
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')

# Plot the regression lines
plot_xs = []
for i in np.arange(N):
    if i == 0:
        plot_x = np.vstack((np.linspace(min(x), max(x),50), np.ones(50)))
    else:
        plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))
    plot_xs.append(plot_x)

for i in np.arange(N):
    ax.plot(plot_xs[i][-2,:], thetas[i].dot(plot_xs[i]))

labels = ['data']
[labels.append('n={}'.format(i+1)) for i in np.arange(N)]
bbox_to_anchor=(1.3, 1)
lgd = ax.legend(labels, bbox_to_anchor=bbox_to_anchor)
```



Calculating the training error (10 points)

Here, we'll now calculate the training error of polynomial models of orders 1 to 5:

$$L(\theta) = \frac{1}{2} \sum_j (\hat{y}_j - y_j)^2$$

In [8]:

```
training_errors = []

# ===== #
# START YOUR CODE HERE #
# ===== #

# GOAL: create a variable training_errors, a list of 5 elements,
# where training_errors[i] are the training loss for the polynomial fit of order i+1.
for i in range(N):
    training_errors.append(np.sum((y - thetas[i].dot(xhats[i]))**2/2))
# ===== #
# END YOUR CODE HERE #
# ===== #

print ('Training errors are: \n', training_errors)
```

Training errors are:

```
[0.2379961088362701, 0.10924922209268531, 0.08169603801105374, 0.08165353735296982, 0.08161479195525298]
```

QUESTIONS

(1) Which polynomial model has the best training error?

(2) Why is this expected?

ANSWERS

(1) Polynomial model of **order 5** has the best training error.

(2) This is expected because polynomial of higher degree will try to **overfit** the data.

Generating new samples and validation error (5 points)


Here, we'll now generate new samples and calculate the validation error of polynomial models of orders 1 to 5.

In [9]:

```
x = np.random.uniform(low=1, high=2, size=(num_train,))
y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')
```

Out[9]:

```
Text(0, 0.5, '$y$')
```



```
In [10]:
xhats = []
for i in np.arange(N):
    if i == 0:
        xhat = np.vstack((x, np.ones_like(x)))
        plot_x = np.vstack((np.linspace(min(x), max(x), 50), np.ones(50)))
    else:
        xhat = np.vstack((x**(i+1), xhat))
        plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))


    xhats.append(xhat)
```

```
In [11]:
# Plot the data
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')

# Plot the regression lines
plot_xs = []
for i in np.arange(N):
    if i == 0:
        plot_x = np.vstack((np.linspace(min(x), max(x), 50), np.ones(50)))
    else:
        plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))
    plot_xs.append(plot_x)

for i in np.arange(N):
    ax.plot(plot_xs[i][-2:], thetas[i].dot(plot_xs[i]))

labels = ['data']
[labels.append('n={}'.format(i+1)) for i in np.arange(N)]
bbox_to_anchor=(1.3, 1)
lgd = ax.legend(labels, bbox_to_anchor=bbox_to_anchor)
```



```
In [ ]: validation_errors = []

# ===== #
# START YOUR CODE HERE #
# ===== #

# GOAL: create a variable validation_errors, a list of 5 elements,
# where validation_errors[i] are the validation loss for the polynomial fit of order i+1.
for i in range(N):
    validation_errors.append(np.sum((y - thetas[i].dot(xhats[i]))**2/2))

# ===== #
# END YOUR CODE HERE #
# ===== #

print ('Validation errors are: \n', validation_errors)
```

QUESTIONS

- (1) Which polynomial model has the best validation error?
- (2) Why does the order-5 polynomial model not generalize well?

ANSWERS

- (1) Polynomial model of **order 4** has the best validation error.
- (2) order-5 polynomial model didn't generalize well because it overfitted the training data in range 0 to 1, so when it was given data outside this range, it gave high errors.