

Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

```
In [2]: ## Import and setups

import time

import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from utils.data_utils import get_CIFAR10_data
from utils.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from utils.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses `for` loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple `for` loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

```
In [17]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [ 0.21027089,  0.21661097],
                          [ 0.22847626,  0.23004637]],
                         [[ 0.50813986,  0.54309974],
```

```

[ 0.64082444, 0.67101435]],
[[[-0.98053589, -1.03143541],
 [-1.19128892, -1.24695841]],
[[ 0.69108355, 0.66880383],
 [ 0.59480972, 0.56776003]],
[[ 2.36270298, 2.36904306],
 [ 2.38090835, 2.38247847]]]])

```

```

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))

```

```

Testing conv_forward_naive
difference: 2.2121476575931688e-08

```

Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple `for` loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

In [31]:

```

x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_forward_naive(x,w,b,conv_param)

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b, conv_param)[0], x,
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b, conv_param)[0], w,
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b, conv_param)[0], b,

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9'
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))

```

```

(4, 3, 5, 5) (4, 2, 5, 5)
Testing conv_backward_naive function
dx error: 2.8496324342528193e-08
dw error: 1.145905092285234e-09
db error: 3.796437261992522e-12

```

Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

In [37]:

```

x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],

```

```

        [-0.20421053, -0.18947368]],
        [[-0.14526316, -0.13052632],
         [-0.08631579, -0.07157895]],
        [[-0.02736842, -0.01263158],
         [ 0.03157895,  0.04631579]]],
        [[[ 0.09052632,  0.10526316],
          [ 0.14947368,  0.16421053]],
         [[ 0.20842105,  0.22315789],
          [ 0.26736842,  0.28210526]],
         [[ 0.32631579,  0.34105263],
          [ 0.38526316,  0.4          ]]]])

```

```

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

```

```

Testing max_pool_forward_naive function:
difference:  4.1666665157267834e-08

```

Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

In [41]:

```

x = np.random.randn(3, 2, 8, 8)
dout = np.random.randn(3, 2, 4, 4)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x, pool_param)[0], x, d

out, cache = max_pool_forward_naive(x, pool_param)
dx = max_pool_backward_naive(dout, cache)

# Your error should be around 1e-12
print('Testing max_pool_backward_naive function:')
print('dx error: ', rel_error(dx, dx_num))

```

```

Testing max_pool_backward_naive function:
dx error:  3.2756245539905995e-12

```

Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by utils. They are provided in `utils/fast_layers.py`.

The fast convolution implementation depends on a Cython extension ('pip install Cython' to your virtual environment); to compile it you need to run the following from the `utils` directory:

```
python setup.py build_ext --inplace
```

NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

```
In [42]: from utils.fast_layers import conv_forward_fast, conv_backward_fast
from time import time

x = np.random.randn(100, 3, 31, 31)
w = np.random.randn(25, 3, 3, 3)
b = np.random.randn(25,)
dout = np.random.randn(100, 25, 16, 16)
conv_param = {'stride': 2, 'pad': 1}

t0 = time()
out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
t1 = time()
out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
t2 = time()

print('Testing conv_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('Difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
t1 = time()
dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting conv_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
print('dw difference: ', rel_error(dw_naive, dw_fast))
print('db difference: ', rel_error(db_naive, db_fast))

Testing conv_forward_fast:
Naive: 14.657998s
Fast: 0.479885s
Speedup: 30.544792x
Difference: 4.3381416333149016e-12

Testing conv_backward_fast:
Naive: 37.213266s
Fast: 0.009992s
Speedup: 3724.438055x
dx difference: 1.9797083570797586e-11
dw difference: 8.705056450802472e-13
db difference: 0.0

In [43]: from utils.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
```

```

dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))

```

```

Testing pool_forward_fast:
Naive: 0.388547s
fast: 0.003029s
speedup: 128.271232x
difference: 0.0

```

```

Testing pool_backward_fast:
Naive: 0.691170s
speedup: 63.002065x
dx difference: 0.0

```

Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py` :

- `conv_relu_forward`
- `conv_relu_backward`
- `conv_relu_pool_forward`
- `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

In [46]:

```

from nndl.conv_layer_utils import conv_relu_pool_forward, conv_relu_pool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w, b, conv_param, po
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w, b, conv_param, po
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w, b, conv_param, po

print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

```

Testing conv_relu_pool
dx error: 4.5124399231583796e-08
dw error: 6.170788280844852e-10
db error: 2.2661151213160115e-11

```

In [47]:

```

from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

```

```
dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b, conv_param)[0], x,  
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b, conv_param)[0], w,  
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b, conv_param)[0], b,  
  
print('Testing conv_relu:')  
print('dx error: ', rel_error(dx_num, dx))  
print('dw error: ', rel_error(dw_num, dw))  
print('db error: ', rel_error(db_num, db))
```

Testing conv_relu:

dx error: 1.8323996630127003e-09

dw error: 1.0764646067955963e-09

db error: 4.147675486357139e-11

What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape (N, D) and produces outputs of shape (N, D) , where we normalize across the minibatch dimension N . For data coming from convolutional layers, batch normalization accepts inputs of shape (N, C, H, W) and produces outputs of shape (N, C, H, W) where the N dimension gives the minibatch size and the (H, W) dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the C feature maps we have (i.e., the layer has C filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the (N, C, H, W) array as an $(N * H * W, C)$ array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- layers.py for your FC network layers, as well as batchnorm and dropout.
- layer_utils.py for your combined FC network layers.
- optim.py for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the batchnorm, then your spatial batchnorm implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

```
In [1]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from utils.data_utils import get_CIFAR10_data
from utils.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from utils.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
```

```
""" returns relative error """
return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
In [15]: # Check the training-time forward pass by checking means and variances
# of features both before and after spatial batch normalization

N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print('  Shape: ', x.shape)
print('  Means: ', x.mean(axis=(0, 2, 3)))
print('  Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma
gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))
```

```
Before spatial batch normalization:
  Shape: (2, 3, 4, 5)
  Means: [ 9.24079377 11.41428531 10.43869742]
  Stds: [3.42195716 3.50733199 3.75066663]
After spatial batch normalization:
  Shape: (2, 3, 4, 5)
  Means: [-4.44089210e-16  8.27116153e-16  5.66213743e-16]
  Stds: [0.99999957 0.99999959 0.99999964]
After spatial batch normalization (nontrivial gamma, beta):
  Shape: (2, 3, 4, 5)
  Means: [6. 7. 8.]
  Stds: [2.99999872 3.99999837 4.99999822]
```

Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
In [18]: N, C, H, W = 2, 3, 4, 5
x = 5 * np.random.randn(N, C, H, W) + 12
gamma = np.random.randn(C)
beta = np.random.randn(C)
dout = np.random.randn(N, C, H, W)

bn_param = {'mode': 'train'}
fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
```



```
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)
dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))
```

```
dx error:  5.68794263178987e-09
dgamma error:  3.842018013168464e-12
dbeta error:  3.275106490366651e-12
```

In []:

Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve > 65% validation error on CIFAR-10.

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- layers.py for your FC network layers, as well as batchnorm and dropout.
- layer_utils.py for your combined FC network layers.
- optim.py for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

```
In [1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from utils.data_utils import get_CIFAR10_data
from utils.gradient_check import eval_numerical_gradient_array, eval_numerical_gradient
from nndl.layers import *
from nndl.conv_layers import *
from utils.fast_layers import *
from utils.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
In [4]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))

X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `nnd1/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your `W1` max relative error and `W2` max relative error are around or below 0.01, they should be acceptable. Other errors should be less than $1e-5$.

```
In [8]: num_inputs = 2
input_dim = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_dim)
y = np.random.randint(num_classes, size=num_inputs)

model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                           input_dim=input_dim, hidden_dim=7,
                           dtype=np.float64)
loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name], verbose=False, h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_

W1 max relative error: 0.011582167206514055
W2 max relative error: 0.014215630010808635
W3 max relative error: 7.143934702505574e-05
b1 max relative error: 1.8072548730142395e-05
b2 max relative error: 1.5045429753527004e-06
b3 max relative error: 1.0474466450818472e-09
```

Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.

```
In [9]: num_train = 100
small_data = {
    'X_train': data['X_train'][:num_train],
    'y_train': data['y_train'][:num_train],
    'X_val': data['X_val'],
    'y_val': data['y_val'],
}

model = ThreeLayerConvNet(weight_scale=1e-2)

solver = Solver(model, small_data,
                 num_epochs=10, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=1)
solver.train()
```


(Iteration 1 / 20) loss: 2.382580

```
(Epoch 0 / 10) train acc: 0.230000; val_acc: 0.127000
(Iteration 2 / 20) loss: 3.858268
(Epoch 1 / 10) train acc: 0.200000; val_acc: 0.114000
(Iteration 3 / 20) loss: 3.676738
(Iteration 4 / 20) loss: 2.294921
(Epoch 2 / 10) train acc: 0.250000; val_acc: 0.108000
(Iteration 5 / 20) loss: 2.364957
(Iteration 6 / 20) loss: 2.177080
(Epoch 3 / 10) train acc: 0.190000; val_acc: 0.128000
(Iteration 7 / 20) loss: 2.091676
(Iteration 8 / 20) loss: 1.936924
(Epoch 4 / 10) train acc: 0.350000; val_acc: 0.152000
(Iteration 9 / 20) loss: 1.948134
(Iteration 10 / 20) loss: 1.703696
(Epoch 5 / 10) train acc: 0.530000; val_acc: 0.175000
(Iteration 11 / 20) loss: 1.607247
(Iteration 12 / 20) loss: 1.332345
(Epoch 6 / 10) train acc: 0.600000; val_acc: 0.182000
(Iteration 13 / 20) loss: 1.358965
(Iteration 14 / 20) loss: 1.301250
(Epoch 7 / 10) train acc: 0.710000; val_acc: 0.217000
(Iteration 15 / 20) loss: 1.181836
(Iteration 16 / 20) loss: 1.537372
(Epoch 8 / 10) train acc: 0.730000; val_acc: 0.214000
(Iteration 17 / 20) loss: 0.771859
(Iteration 18 / 20) loss: 0.951309
(Epoch 9 / 10) train acc: 0.730000; val_acc: 0.198000
(Iteration 19 / 20) loss: 0.712975
(Iteration 20 / 20) loss: 0.561128
(Epoch 10 / 10) train acc: 0.820000; val_acc: 0.209000
```

In [10]:

```
plt.subplot(2, 1, 1)
plt.plot(solver.loss_history, 'o')
plt.xlabel('iteration')
plt.ylabel('loss')

plt.subplot(2, 1, 2)
plt.plot(solver.train_acc_history, '-o')
plt.plot(solver.val_acc_history, '-o')
plt.legend(['train', 'val'], loc='upper left')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.show()
```



Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

```
In [11]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=1, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)

solver.train()
```

```
(Iteration 1 / 980) loss: 2.304531
(Epoch 0 / 1) train acc: 0.096000; val_acc: 0.098000
(Iteration 21 / 980) loss: 2.271377
(Iteration 41 / 980) loss: 2.002693
(Iteration 61 / 980) loss: 1.996322
(Iteration 81 / 980) loss: 1.863851
(Iteration 101 / 980) loss: 2.064730
(Iteration 121 / 980) loss: 1.992802
(Iteration 141 / 980) loss: 1.452346
(Iteration 161 / 980) loss: 1.869969
(Iteration 181 / 980) loss: 1.793654
(Iteration 201 / 980) loss: 2.008249
(Iteration 221 / 980) loss: 1.715329
(Iteration 241 / 980) loss: 1.515369
(Iteration 261 / 980) loss: 1.675155
(Iteration 281 / 980) loss: 1.806337
(Iteration 301 / 980) loss: 1.473572
(Iteration 321 / 980) loss: 1.447883
(Iteration 341 / 980) loss: 1.981338
(Iteration 361 / 980) loss: 1.718061
(Iteration 381 / 980) loss: 1.418680
(Iteration 401 / 980) loss: 1.740535
(Iteration 421 / 980) loss: 1.706089
```



```
(Iteration 441 / 980) loss: 1.512991
(Iteration 461 / 980) loss: 1.668131
(Iteration 481 / 980) loss: 1.629799
(Iteration 501 / 980) loss: 1.625585
(Iteration 521 / 980) loss: 1.501371
(Iteration 541 / 980) loss: 1.379852
(Iteration 561 / 980) loss: 1.410839
(Iteration 581 / 980) loss: 1.368521
(Iteration 601 / 980) loss: 1.558027
(Iteration 621 / 980) loss: 1.799553
(Iteration 641 / 980) loss: 1.564854
(Iteration 661 / 980) loss: 1.369897
(Iteration 681 / 980) loss: 1.381831
(Iteration 701 / 980) loss: 1.705887
(Iteration 721 / 980) loss: 1.601044
(Iteration 741 / 980) loss: 1.451873
(Iteration 761 / 980) loss: 1.671530
(Iteration 781 / 980) loss: 1.789960
(Iteration 801 / 980) loss: 1.541475
(Iteration 821 / 980) loss: 1.907693
(Iteration 841 / 980) loss: 1.492666
(Iteration 861 / 980) loss: 1.340623
(Iteration 881 / 980) loss: 1.759223
(Iteration 901 / 980) loss: 1.692885
(Iteration 921 / 980) loss: 1.639575
(Iteration 941 / 980) loss: 1.656631
(Iteration 961 / 980) loss: 1.858810
(Epoch 1 / 1) train acc: 0.462000; val_acc: 0.433000
```

Get > 65% validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]XN - [affine]XM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

```
In [13]: # ===== #
# YOUR CODE HERE:
# Implement a CNN to achieve greater than 65% validation accuracy
# on CIFAR-10.
```

```
# ===== #
# tuning hyperparameters achieves >65% validation accuracy
model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=10, batch_size=200,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 5e-4,
                 },
                 lr_decay=0.95,
                 verbose=True, print_every=100)
solver.train()

# ===== #
# END YOUR CODE HERE
# ===== #
```

```
(Iteration 1 / 2450) loss: 2.304694
(Epoch 0 / 10) train acc: 0.085000; val_acc: 0.079000
(Iteration 101 / 2450) loss: 1.515280
(Iteration 201 / 2450) loss: 1.260570
(Epoch 1 / 10) train acc: 0.581000; val_acc: 0.555000
(Iteration 301 / 2450) loss: 1.188566
(Iteration 401 / 2450) loss: 1.167816
(Epoch 2 / 10) train acc: 0.644000; val_acc: 0.604000
(Iteration 501 / 2450) loss: 0.943109
(Iteration 601 / 2450) loss: 0.989317
(Iteration 701 / 2450) loss: 0.836393
(Epoch 3 / 10) train acc: 0.685000; val_acc: 0.614000
(Iteration 801 / 2450) loss: 0.987970
(Iteration 901 / 2450) loss: 1.054005
(Epoch 4 / 10) train acc: 0.727000; val_acc: 0.636000
(Iteration 1001 / 2450) loss: 0.963190
(Iteration 1101 / 2450) loss: 0.841740
(Iteration 1201 / 2450) loss: 0.807908
(Epoch 5 / 10) train acc: 0.765000; val_acc: 0.657000
(Iteration 1301 / 2450) loss: 0.740925
(Iteration 1401 / 2450) loss: 0.644606
(Epoch 6 / 10) train acc: 0.778000; val_acc: 0.661000
(Iteration 1501 / 2450) loss: 0.623828
(Iteration 1601 / 2450) loss: 0.649584
(Iteration 1701 / 2450) loss: 0.727654
(Epoch 7 / 10) train acc: 0.818000; val_acc: 0.646000
(Iteration 1801 / 2450) loss: 0.611627
(Iteration 1901 / 2450) loss: 0.639321
(Epoch 8 / 10) train acc: 0.829000; val_acc: 0.665000
(Iteration 2001 / 2450) loss: 0.581038
(Iteration 2101 / 2450) loss: 0.390481
(Iteration 2201 / 2450) loss: 0.517170
(Epoch 9 / 10) train acc: 0.878000; val_acc: 0.654000
(Iteration 2301 / 2450) loss: 0.501238
(Iteration 2401 / 2450) loss: 0.476331
(Epoch 10 / 10) train acc: 0.905000; val_acc: 0.671000
```

In []:

```

1 import numpy as np
2
3 from nn dl.layers import *
4 from nn dl.conv_layers import *
5 from utils.fast_layers import *
6 from nn dl.layer_utils import *
7 from nn dl.conv_layer_utils import *
8
9 import pdb
10
11 class ThreeLayerConvNet(object):
12     """
13     A three-layer convolutional network with the following architecture:
14
15     conv - relu - 2x2 max pool - affine - relu - affine - softmax
16
17     The network operates on minibatches of data that have shape (N, C, H, W)
18     consisting of N images, each with height H and width W and with C input
19     channels.
20     """
21
22     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
23                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
24                 dtype=np.float32, use_batchnorm=False):
25         """
26         Initialize a new network.
27
28         Inputs:
29         - input_dim: Tuple (C, H, W) giving size of input data
30         - num_filters: Number of filters to use in the convolutional layer
31         - filter_size: Size of filters to use in the convolutional layer
32         - hidden_dim: Number of units to use in the fully-connected hidden layer
33         - num_classes: Number of scores to produce from the final affine layer.
34         - weight_scale: Scalar giving standard deviation for random initialization
35           of weights.
36         - reg: Scalar giving L2 regularization strength
37         - dtype: numpy datatype to use for computation.
38         """
39         self.use_batchnorm = use_batchnorm
40         self.params = {}
41         self.reg = reg
42         self.dtype = dtype
43
44
45         # ===== #
46         # YOUR CODE HERE:
47         #   Initialize the weights and biases of a three layer CNN. To initialize:
48         #   - the biases should be initialized to zeros.
49         #   - the weights should be initialized to a matrix with entries
50         #     drawn from a Gaussian distribution with zero mean and
51         #     standard deviation given by weight_scale.
52         # ===== #
53
54         self.params['W1'] = np.random.normal(loc=0.0, scale=weight_scale,
55                                             size=(num_filters,input_dim[0],filter_size,filter_size))
56         self.params['b1'] = np.zeros((num_filters))
57         self.params['W2'] = np.random.normal(loc=0.0, scale=weight_scale,
58                                             size=(int(num_filters*input_dim[1]*input_dim[2]/4),hidden_dim))
59         self.params['b2'] = np.zeros((hidden_dim))
60         self.params['W3'] = np.random.normal(loc=0.0, scale=weight_scale,
61                                             size=(hidden_dim,num_classes))
62         self.params['b3'] = np.zeros((num_classes))
63
64         # ===== #
65         # END YOUR CODE HERE
66         # ===== #
67
68         for k, v in self.params.items():
69             self.params[k] = v.astype(dtype)
70
71
72     def loss(self, X, y=None):
73         """

```

Evaluate loss and gradient for the three-layer convolutional network.

Input / output: Same API as TwoLayerNet in fc_net.py.

```
"""
W1, b1 = self.params['W1'], self.params['b1']
W2, b2 = self.params['W2'], self.params['b2']
W3, b3 = self.params['W3'], self.params['b3']

# pass conv_param to the forward pass for the convolutional layer
filter_size = W1.shape[2]
conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}

# pass pool_param to the forward pass for the max-pooling layer
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

scores = None

# ===== #
# YOUR CODE HERE:
# Implement the forward pass of the three layer CNN. Store the output
# scores as the variable "scores".
# ===== #
a,c1 = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
a,c2 = affine_relu_forward(a, W2, b2)
scores,c3 = affine_forward(a,W3,b3)
# ===== #
# END YOUR CODE HERE
# ===== #

if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backward pass of the three layer CNN. Store the grads
# in the grads dictionary, exactly as before (i.e., the gradient of
# self.params[k] will be grads[k]). Store the loss as "loss", and
# don't forget to add regularization on ALL weight matrices.
# ===== #

loss, ds = softmax_loss(scores,y)
da,grads['W3'],grads['b3'] = affine_backward(ds,c3)
da,grads['W2'],grads['b2'] = affine_relu_backward(da,c2)
da,grads['W1'],grads['b1'] = conv_relu_pool_backward(da,c1)
for w in ['W1','W2','W3']:
    grads[w] += self.reg*self.params[w]
    loss += self.reg*0.5*np.sum(np.square(self.params[w]))
# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads
```

pass

```

1  import numpy as np
2  from nn.layers import *
3  import pdb
4
5
6  def conv_forward_naive(x, w, b, conv_param):
7      """
8      A naive implementation of the forward pass for a convolutional layer.
9
10     The input consists of N data points, each with C channels, height H and width
11     W. We convolve each input with F different filters, where each filter spans
12     all C channels and has height HH and width WW.
13
14     Input:
15     - x: Input data of shape (N, C, H, W)
16     - w: Filter weights of shape (F, C, HH, WW)
17     - b: Biases, of shape (F,)
18     - conv_param: A dictionary with the following keys:
19         - 'stride': The number of pixels between adjacent receptive fields in the
20           horizontal and vertical directions.
21         - 'pad': The number of pixels that will be used to zero-pad the input.
22
23     Returns a tuple of:
24     - out: Output data, of shape (N, F, H', W') where H' and W' are given by
25         H' = 1 + (H + 2 * pad - HH) / stride
26         W' = 1 + (W + 2 * pad - WW) / stride
27     - cache: (x, w, b, conv_param)
28     """
29     out = None
30     pad = conv_param['pad']
31     stride = conv_param['stride']
32
33     # ===== #
34     # YOUR CODE HERE:
35     # Implement the forward pass of a convolutional neural network.
36     # Store the output as 'out'.
37     # Hint: to pad the array, you can use the function np.pad.
38     # ===== #
39     N,C,H,W = x.shape
40     F, C, HH, WW = w.shape
41     npad = ((0, 0), (0, 0), (pad, pad), (pad, pad))
42     xpad = np.pad(x,npad)
43
44     H_ = int(1 + (H + 2 * pad - HH) / stride)
45     W_ = int(1 + (W + 2 * pad - WW) / stride)
46     out = np.zeros((N,F,H_,W_))
47     for i in range(N):
48         for j in range(F):
49             for h in range(H_):
50                 for width in range(W_):
51                     cur=0
52                     for i2 in range(C):
53                         for j2 in range(HH):
54                             for k2 in range(WW):
55                                 cur+=xpad[i,i2,h*stride+j2,width*stride+k2]*w[j,i2,j2,k2]
56                             out[i,j,h,width] = cur + b[j]
57
58     # ===== #
59     # END YOUR CODE HERE
60     # ===== #
61
62     cache = (x, w, b, conv_param)
63     return out, cache
64
65 def conv_backward_naive(dout, cache):
66     """
67     A naive implementation of the backward pass for a convolutional layer.
68
69     Inputs:
70     - dout: Upstream derivatives.
71     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
72

```



```

73 Returns a tuple of:
74 - dx: Gradient with respect to x
75 - dw: Gradient with respect to w
76 - db: Gradient with respect to b
77 """
78 dx, dw, db = None, None, None
79
80 N, F, out_height, out_width = dout.shape
81 x, w, b, conv_param = cache
82
83 stride, pad = [conv_param['stride'], conv_param['pad']]
84 xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
85 num_filts, _, f_height, f_width = w.shape
86
87 # ===== #
88 # YOUR CODE HERE:
89 # Implement the backward pass of a convolutional neural network.
90 # Calculate the gradients: dx, dw, and db.
91 # ===== #
92 db=np.sum(dout,axis=(0,2,3))
93 dxpad=np.zeros(xpad.shape)
94 dw=np.zeros(w.shape)
95
96
97 N,C,H,W = x.shape
98 F, C, HH, WW = w.shape
99 H_ = int(1 + (H + 2 * pad - HH) / stride)
100 W_ = int(1 + (W + 2 * pad - WW) / stride)
101 for i in range(N):
102     for j in range(F):
103         for h in range(H_):
104             for width in range(W_):
105                 cur=0
106                 for i2 in range(C):
107                     for j2 in range(HH):
108                         for k2 in range(WW):
109                             #cur+=xpad[i,i2,h*stride+j2,width*stride+k2]*w[j,i2,j2,k2]
110                             dxpad[i,i2,h*stride+j2,width*stride+k2] +=
111                             w[j,i2,j2,k2]*dout[i,j,h,width]
112                             dw[j,i2,j2,k2] +=
113                             xpad[i,i2,h*stride+j2,width*stride+k2]*dout[i,j,h,width]
114                             #out[i,j,h,width] = cur + b[j]
115 dx = dxpad[:, :, pad:-pad, pad:-pad]
116 # ===== #
117 # END YOUR CODE HERE
118 # ===== #
119
120 return dx, dw, db
121
122 def max_pool_forward_naive(x, pool_param):
123     """
124     A naive implementation of the forward pass for a max pooling layer.
125
126     Inputs:
127     - x: Input data, of shape (N, C, H, W)
128     - pool_param: dictionary with the following keys:
129         - 'pool_height': The height of each pooling region
130         - 'pool_width': The width of each pooling region
131         - 'stride': The distance between adjacent pooling regions
132
133     Returns a tuple of:
134     - out: Output data
135     - cache: (x, pool_param)
136     """
137     out = None
138     # ===== #
139     # YOUR CODE HERE:
140     # Implement the max pooling forward pass.
141     # ===== #
142     N, C, H, W = x.shape

```

```

143 ph = pool_param['pool_height']
144 pw = pool_param['pool_width']
145 s = pool_param['stride']
146 H_ = int(1 + (H-ph)/s)
147 W_ = int(1 + (W-pw)/s)
148 out = np.zeros((N,C,H_,W_))
149 for i in range(N):
150     for j in range(C):
151         for h in range(H_):
152             for w in range(W_):
153                 out[i,j,h,w]=np.max(x[i,j,h*s:h*s+ph,w*s:w*s+pw])
154 # ===== #
155 # END YOUR CODE HERE
156 # ===== #
157 cache = (x, pool_param)
158 return out, cache
159
160 def max_pool_backward_naive(dout, cache):
161     """
162     A naive implementation of the backward pass for a max pooling layer.
163
164     Inputs:
165     - dout: Upstream derivatives
166     - cache: A tuple of (x, pool_param) as in the forward pass.
167
168     Returns:
169     - dx: Gradient with respect to x
170     """
171     dx = None
172     x, pool_param = cache
173     pool_height, pool_width, stride = pool_param['pool_height'],
174     pool_param['pool_width'], pool_param['stride']
175 # ===== #
176 # YOUR CODE HERE:
177 # Implement the max pooling backward pass.
178 # ===== #
179 dx = np.zeros(x.shape)
180 N, C, H, W = x.shape
181 ph = pool_param['pool_height']
182 pw = pool_param['pool_width']
183 s = pool_param['stride']
184 H_ = int(1 + (H-ph)/s)
185 W_ = int(1 + (W-pw)/s)
186 for i in range(N):
187     for j in range(C):
188         for h in range(H_):
189             for w in range(W_):
190                 cur=np.max(x[i,j,h*s:h*s+ph,w*s:w*s+pw])
191                 for i2 in range(ph):
192                     for j2 in range(pw):
193                         if x[i,j,h*s+i2,w*s+j2]==cur:
194                             dx[i,j,h*s+i2,w*s+j2]=dout[i,j,h,w]
195 # ===== #
196 # END YOUR CODE HERE
197 # ===== #
198
199 return dx
200
201
202 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
203     """
204     Computes the forward pass for spatial batch normalization.
205
206     Inputs:
207     - x: Input data of shape (N, C, H, W)
208     - gamma: Scale parameter, of shape (C,)
209     - beta: Shift parameter, of shape (C,)
210     - bn_param: Dictionary with the following keys:
211         - mode: 'train' or 'test'; required
212         - eps: Constant for numeric stability
213         - momentum: Constant for running mean / variance. momentum=0 means that

```

```

214         old information is discarded completely at every time step, while
215         momentum=1 means that new information is never incorporated. The
216         default of momentum=0.9 should work well in most situations.
217     - running_mean: Array of shape (D,) giving running mean of features
218     - running_var: Array of shape (D,) giving running variance of features
219
220 Returns a tuple of:
221 - out: Output data, of shape (N, C, H, W)
222 - cache: Values needed for the backward pass
223 """
224 out, cache = None, None
225
226 # ===== #
227 # YOUR CODE HERE:
228 #     Implement the spatial batchnorm forward pass.
229 #
230 #     You may find it useful to use the batchnorm forward pass you
231 #     implemented in HW #4.
232 # ===== #
233 N, C, H, W = x.shape
234 xb = np.transpose(x, (0,2,3,1))
235 xb = xb.reshape((-1,C))
236 out,cache = batchnorm_forward(xb, gamma, beta, bn_param)
237 out = out.reshape((N,H,W,C))
238 out = np.transpose(out, (0,3,1,2))
239 # ===== #
240 # END YOUR CODE HERE
241 # ===== #
242
243 return out, cache
244
245
246 def spatial_batchnorm_backward(dout, cache):
247     """
248     Computes the backward pass for spatial batch normalization.
249
250     Inputs:
251     - dout: Upstream derivatives, of shape (N, C, H, W)
252     - cache: Values from the forward pass
253
254     Returns a tuple of:
255     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
256     - dgamma: Gradient with respect to scale parameter, of shape (C,)
257     - dbeta: Gradient with respect to shift parameter, of shape (C,)
258     """
259     dx, dgamma, dbeta = None, None, None
260
261     # ===== #
262     # YOUR CODE HERE:
263     #     Implement the spatial batchnorm backward pass.
264     #
265     #     You may find it useful to use the batchnorm forward pass you
266     #     implemented in HW #4.
267     # ===== #
268     N, C, H, W = dout.shape
269     dout = np.transpose(dout, (0,2,3,1))
270     dout = dout.reshape((-1,C))
271     #x,x_n,mean,var,eps,gamma = cache
272     #x = x.reshape((N,H,W,C))
273     #x = np.transpose(x, (0,3,1,2))
274     #x_n = x_n.reshape((N,H,W,C))
275     #x_n = np.transpose(x_n, (0,3,1,2))
276     #cache = x,x_n,mean,var,eps,gamma
277     dx, dgamma, dbeta = batchnorm_backward(dout,cache)
278     dx = dx.reshape((N,H,W,C))
279     dx = np.transpose(dx, (0,3,1,2))
280     # ===== #
281     # END YOUR CODE HERE
282     # ===== #
283
284     return dx, dgamma, dbeta

```

```

1  from nndl.layers import *
2  from utils.fast_layers import *
3
4
5  def conv_relu_forward(x, w, b, conv_param):
6      """
7      A convenience layer that performs a convolution followed by a ReLU.
8
9      Inputs:
10     - x: Input to the convolutional layer
11     - w, b, conv_param: Weights and parameters for the convolutional layer
12
13     Returns a tuple of:
14     - out: Output from the ReLU
15     - cache: Object to give to the backward pass
16     """
17     a, conv_cache = conv_forward_fast(x, w, b, conv_param)
18     out, relu_cache = relu_forward(a)
19     cache = (conv_cache, relu_cache)
20     return out, cache
21
22
23  def conv_relu_backward(dout, cache):
24      """
25      Backward pass for the conv-relu convenience layer.
26      """
27     conv_cache, relu_cache = cache
28     da = relu_backward(dout, relu_cache)
29     dx, dw, db = conv_backward_fast(da, conv_cache)
30     return dx, dw, db
31
32
33  def conv_relu_pool_forward(x, w, b, conv_param, pool_param):
34      """
35      Convenience layer that performs a convolution, a ReLU, and a pool.
36
37      Inputs:
38     - x: Input to the convolutional layer
39     - w, b, conv_param: Weights and parameters for the convolutional layer
40     - pool_param: Parameters for the pooling layer
41
42     Returns a tuple of:
43     - out: Output from the pooling layer
44     - cache: Object to give to the backward pass
45     """
46     a, conv_cache = conv_forward_fast(x, w, b, conv_param)
47     s, relu_cache = relu_forward(a)
48     out, pool_cache = max_pool_forward_fast(s, pool_param)
49     cache = (conv_cache, relu_cache, pool_cache)
50     return out, cache
51
52
53  def conv_relu_pool_backward(dout, cache):
54      """
55      Backward pass for the conv-relu-pool convenience layer
56      """
57     conv_cache, relu_cache, pool_cache = cache
58     ds = max_pool_backward_fast(dout, pool_cache)
59     da = relu_backward(ds, relu_cache)
60     dx, dw, db = conv_backward_fast(da, conv_cache)
61     return dx, dw, db

```

```

1  import numpy as np
2  import pdb
3
4  def affine_forward(x, w, b):
5      """
6      Computes the forward pass for an affine (fully-connected) layer.
7
8      The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
9      examples, where each example x[i] has shape (d_1, ..., d_k). We will
10     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
11     then transform it to an output vector of dimension M.
12
13     Inputs:
14     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
15     - w: A numpy array of weights, of shape (D, M)
16     - b: A numpy array of biases, of shape (M,)
17
18     Returns a tuple of:
19     - out: output, of shape (N, M)
20     - cache: (x, w, b)
21     """
22
23     # ===== #
24     # YOUR CODE HERE:
25     # Calculate the output of the forward pass. Notice the dimensions
26     # of w are D x M, which is the transpose of what we did in earlier
27     # assignments.
28     # ===== #
29
30     out = x.reshape(x.shape[0],-1).dot(w)+b
31     # ===== #
32     # END YOUR CODE HERE
33     # ===== #
34
35     cache = (x, w, b)
36     return out, cache
37
38
39  def affine_backward(dout, cache):
40      """
41      Computes the backward pass for an affine layer.
42
43      Inputs:
44      - dout: Upstream derivative, of shape (N, M)
45      - cache: Tuple of:
46        - x: Input data, of shape (N, d_1, ... d_k)
47        - w: Weights, of shape (D, M)
48
49      Returns a tuple of:
50      - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
51      - dw: Gradient with respect to w, of shape (D, M)
52      - db: Gradient with respect to b, of shape (M,)
53      """
54      x, w, b = cache
55      dx, dw, db = None, None, None
56
57      # ===== #
58      # YOUR CODE HERE:
59      # Calculate the gradients for the backward pass.
60      # ===== #
61
62      dx = dout.dot(w.T).reshape(x.shape)
63      dw = x.reshape(x.shape[0],-1).T.dot(dout).reshape(w.shape)
64      db = np.sum(dout,axis=0)
65      # ===== #
66      # END YOUR CODE HERE
67      # ===== #
68
69      return dx, dw, db
70
71  def relu_forward(x):
72      """

```



```

73     Computes the forward pass for a layer of rectified linear units (ReLUs).
74
75     Input:
76     - x: Inputs, of any shape
77
78     Returns a tuple of:
79     - out: Output, of the same shape as x
80     - cache: x
81     """
82     # ===== #
83     # YOUR CODE HERE:
84     #     Implement the ReLU forward pass.
85     # ===== #
86     out = np.maximum(x,0)
87     # ===== #
88     # END YOUR CODE HERE
89     # ===== #
90
91     cache = x
92     return out, cache
93
94
95 def relu_backward(dout, cache):
96     """
97     Computes the backward pass for a layer of rectified linear units (ReLUs).
98
99     Input:
100     - dout: Upstream derivatives, of any shape
101     - cache: Input x, of same shape as dout
102
103     Returns:
104     - dx: Gradient with respect to x
105     """
106     x = cache
107
108     # ===== #
109     # YOUR CODE HERE:
110     #     Implement the ReLU backward pass
111     # ===== #
112
113     dx = dout
114     dx[x<0]=0
115     # ===== #
116     # END YOUR CODE HERE
117     # ===== #
118
119     return dx
120
121 def batchnorm_forward(x, gamma, beta, bn_param):
122     """
123     Forward pass for batch normalization.
124
125     During training the sample mean and (uncorrected) sample variance are
126     computed from minibatch statistics and used to normalize the incoming data.
127     During training we also keep an exponentially decaying running mean of the mean
128     and variance of each feature, and these averages are used to normalize data
129     at test-time.
130
131     At each timestep we update the running averages for mean and variance using
132     an exponential decay based on the momentum parameter:
133
134     running_mean = momentum * running_mean + (1 - momentum) * sample_mean
135     running_var = momentum * running_var + (1 - momentum) * sample_var
136
137     Note that the batch normalization paper suggests a different test-time
138     behavior: they compute sample mean and variance for each feature using a
139     large number of training images rather than using a running average. For
140     this implementation we have chosen to use running averages instead since
141     they do not require an additional estimation step; the torch7 implementation
142     of batch normalization also uses running averages.
143
144     Input:

```

```

145 - x: Data of shape (N, D)
146 - gamma: Scale parameter of shape (D,)
147 - beta: Shift parameter of shape (D,)
148 - bn_param: Dictionary with the following keys:
149   - mode: 'train' or 'test'; required
150   - eps: Constant for numeric stability
151   - momentum: Constant for running mean / variance.
152   - running_mean: Array of shape (D,) giving running mean of features
153   - running_var: Array of shape (D,) giving running variance of features
154
155 Returns a tuple of:
156 - out: of shape (N, D)
157 - cache: A tuple of values needed in the backward pass
158 """
159 mode = bn_param['mode']
160 eps = bn_param.get('eps', 1e-5)
161 momentum = bn_param.get('momentum', 0.9)
162
163 N, D = x.shape
164 running_mean = bn_param.get('running_mean', np.zeros(D, dtype=x.dtype))
165 running_var = bn_param.get('running_var', np.zeros(D, dtype=x.dtype))
166
167 out, cache = None, None
168 if mode == 'train':
169
170     # ===== #
171     # YOUR CODE HERE:
172     #   A few steps here:
173     #   (1) Calculate the running mean and variance of the minibatch.
174     #   (2) Normalize the activations with the running mean and variance.
175     #   (3) Scale and shift the normalized activations. Store this
176     #       as the variable 'out'
177     #   (4) Store any variables you may need for the backward pass in
178     #       the 'cache' variable.
179     # ===== #
180
181     cur_mean = np.mean(x, axis=0)
182     cur_var = np.var(x, axis=0)
183     running_mean = momentum * running_mean + (1 - momentum) * cur_mean
184     running_var = momentum * running_var + (1 - momentum) * cur_var
185     x_n = (x - cur_mean) / np.sqrt(cur_var + eps)
186     out = gamma * x_n + beta
187     cache = x, x_n, cur_mean, cur_var, eps, gamma
188     # ===== #
189     # END YOUR CODE HERE
190     # ===== #
191
192 elif mode == 'test':
193
194     # ===== #
195     # YOUR CODE HERE:
196     #   Calculate the testing time normalized activation. Normalize using
197     #   the running mean and variance, and then scale and shift appropriately.
198     #   Store the output as 'out'.
199     # ===== #
200
201     x = (x - running_mean) / np.sqrt(running_var + eps)
202     out = gamma * x + beta
203     # ===== #
204     # END YOUR CODE HERE
205     # ===== #
206
207 else:
208     raise ValueError('Invalid forward batchnorm mode "%s"' % mode)
209
210 # Store the updated running means back into bn_param
211 bn_param['running_mean'] = running_mean
212 bn_param['running_var'] = running_var
213
214 return out, cache
215
216 def batchnorm_backward(dout, cache):

```

```

217 """
218 Backward pass for batch normalization.
219
220 For this implementation, you should write out a computation graph for
221 batch normalization on paper and propagate gradients backward through
222 intermediate nodes.
223
224 Inputs:
225 - dout: Upstream derivatives, of shape (N, D)
226 - cache: Variable of intermediates from batchnorm_forward.
227
228 Returns a tuple of:
229 - dx: Gradient with respect to inputs x, of shape (N, D)
230 - dgamma: Gradient with respect to scale parameter gamma, of shape (D,)
231 - dbeta: Gradient with respect to shift parameter beta, of shape (D,)
232 """
233 dx, dgamma, dbeta = None, None, None
234
235 # ===== #
236 # YOUR CODE HERE:
237 # Implement the batchnorm backward pass, calculating dx, dgamma, and dbeta.
238 # ===== #
239
240 x, x_n, mean, var, eps, gamma = cache
241 dbeta = np.sum(dout, axis=0)
242 dgamma = np.sum(dout*x_n, axis=0)
243 dxh = dout*gamma
244 std = np.sqrt(var+eps)
245 du = -np.sum((1/std)*dxh, axis=0)
246 dv = -np.sum((1/(2*(std**3)))*(x-mean)*dxh, axis=0)
247 dx = dxh*(1/std) + du/x.shape[0] + (2*(x-mean)/x.shape[0])*dv
248 # ===== #
249 # END YOUR CODE HERE
250 # ===== #
251
252 return dx, dgamma, dbeta
253
254 def dropout_forward(x, dropout_param):
255     """
256     Performs the forward pass for (inverted) dropout.
257
258     Inputs:
259     - x: Input data, of any shape
260     - dropout_param: A dictionary with the following keys:
261         - p: Dropout parameter. We drop each neuron output with probability p.
262         - mode: 'test' or 'train'. If the mode is train, then perform dropout;
263           if the mode is test, then just return the input.
264         - seed: Seed for the random number generator. Passing seed makes this
265           function deterministic, which is needed for gradient checking but not in
266           real networks.
267
268     Outputs:
269     - out: Array of the same shape as x.
270     - cache: A tuple (dropout_param, mask). In training mode, mask is the dropout
271       mask that was used to multiply the input; in test mode, mask is None.
272     """
273     p, mode = dropout_param['p'], dropout_param['mode']
274     if 'seed' in dropout_param:
275         np.random.seed(dropout_param['seed'])
276
277     mask = None
278     out = None
279
280     if mode == 'train':
281         # ===== #
282         # YOUR CODE HERE:
283         # Implement the inverted dropout forward pass during training time.
284         # Store the masked and scaled activations in out, and store the
285         # dropout mask as the variable mask.
286         # ===== #
287
288         mask = (np.random.rand(*x.shape) < (1-p)) / (1-p)

```

```

289     out = x*mask
290     # ===== #
291     # END YOUR CODE HERE
292     # ===== #
293
294     elif mode == 'test':
295
296         # ===== #
297         # YOUR CODE HERE:
298         # Implement the inverted dropout forward pass during test time.
299         # ===== #
300         out=x
301
302         # ===== #
303         # END YOUR CODE HERE
304         # ===== #
305
306         cache = (dropout_param, mask)
307         out = out.astype(x.dtype, copy=False)
308
309         return out, cache
310
311 def dropout_backward(dout, cache):
312     """
313     Perform the backward pass for (inverted) dropout.
314
315     Inputs:
316     - dout: Upstream derivatives, of any shape
317     - cache: (dropout_param, mask) from dropout_forward.
318     """
319     dropout_param, mask = cache
320     mode = dropout_param['mode']
321
322     dx = None
323     if mode == 'train':
324         # ===== #
325         # YOUR CODE HERE:
326         # Implement the inverted dropout backward pass during training time.
327         # ===== #
328         (dropout_param, mask) = cache
329         dx = dout*mask
330
331         # ===== #
332         # END YOUR CODE HERE
333         # ===== #
334     elif mode == 'test':
335         # ===== #
336         # YOUR CODE HERE:
337         # Implement the inverted dropout backward pass during test time.
338         # ===== #
339         dx=dout
340         # ===== #
341         # END YOUR CODE HERE
342         # ===== #
343     return dx
344
345 def svm_loss(x, y):
346     """
347     Computes the loss and gradient using for multiclass SVM classification.
348
349     Inputs:
350     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
351         for the ith input.
352     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
353         0 <= y[i] < C
354
355     Returns a tuple of:
356     - loss: Scalar giving the loss
357     - dx: Gradient of the loss with respect to x
358     """
359     N = x.shape[0]
360     correct_class_scores = x[np.arange(N), y]

```

```

361 margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
362 margins[np.arange(N), y] = 0
363 loss = np.sum(margins) / N
364 num_pos = np.sum(margins > 0, axis=1)
365 dx = np.zeros_like(x)
366 dx[margins > 0] = 1
367 dx[np.arange(N), y] -= num_pos
368 dx /= N
369 return loss, dx
370
371
372 def softmax_loss(x, y):
373     """
374     Computes the loss and gradient for softmax classification.
375
376     Inputs:
377     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
378       for the ith input.
379     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
380       0 <= y[i] < C
381
382     Returns a tuple of:
383     - loss: Scalar giving the loss
384     - dx: Gradient of the loss with respect to x
385     """
386
387     probs = np.exp(x - np.max(x, axis=1, keepdims=True))
388     probs /= np.sum(probs, axis=1, keepdims=True)
389     N = x.shape[0]
390     loss = -np.sum(np.log(probs[np.arange(N), y])) / N
391     dx = probs.copy()
392     dx[np.arange(N), y] -= 1
393     dx /= N
394     return loss, dx
395

```



```

1  import numpy as np
2
3  """
4  This file implements various first-order update rules that are commonly used for
5  training neural networks. Each update rule accepts current weights and the
6  gradient of the loss with respect to those weights and produces the next set of
7  weights. Each update rule has the same interface:
8
9  def update(w, dw, config=None):
10
11  Inputs:
12      - w: A numpy array giving the current weights.
13      - dw: A numpy array of the same shape as w giving the gradient of the
14        loss with respect to w.
15      - config: A dictionary containing hyperparameter values such as learning rate,
16        momentum, etc. If the update rule requires caching values over many
17        iterations, then config will also hold these cached values.
18
19  Returns:
20      - next_w: The next point after the update.
21      - config: The config dictionary to be passed to the next iteration of the
22        update rule.
23
24  NOTE: For most update rules, the default learning rate will probably not perform
25  well; however the default values of the other hyperparameters should work well
26  for a variety of different problems.
27
28  For efficiency, update rules may perform in-place updates, mutating w and
29  setting next_w equal to w.
30  """
31
32
33  def sgd(w, dw, config=None):
34      """
35      Performs vanilla stochastic gradient descent.
36
37      config format:
38      - learning_rate: Scalar learning rate.
39      """
40      if config is None: config = {}
41      config.setdefault('learning_rate', 1e-2)
42
43      w -= config['learning_rate'] * dw
44      return w, config
45
46
47  def sgd_momentum(w, dw, config=None):
48      """
49      Performs stochastic gradient descent with momentum.
50
51      config format:
52      - learning_rate: Scalar learning rate.
53      - momentum: Scalar between 0 and 1 giving the momentum value.
54        Setting momentum = 0 reduces to sgd.
55      - velocity: A numpy array of the same shape as w and dw used to store a moving
56        average of the gradients.
57      """
58      if config is None: config = {}
59      config.setdefault('learning_rate', 1e-2)
60      config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
61      v = config.get('velocity', np.zeros_like(w)) # gets velocity, else sets it to
62        zero.
63
64      # ===== #
65      # YOUR CODE HERE:
66      #   Implement the momentum update formula. Return the updated weights
67      #   as next_w, and the updated velocity as v.
68      # ===== #
69      v = config['momentum'] * v - config['learning_rate'] * dw
70      next_w = w+v
71      # ===== #
72      # END YOUR CODE HERE

```

```

72 # ===== #
73
74 config['velocity'] = v
75
76 return next_w, config
77
78 def sgd_nesterov_momentum(w, dw, config=None):
79     """
80     Performs stochastic gradient descent with Nesterov momentum.
81
82     config format:
83     - learning_rate: Scalar learning rate.
84     - momentum: Scalar between 0 and 1 giving the momentum value.
85       Setting momentum = 0 reduces to sgd.
86     - velocity: A numpy array of the same shape as w and dw used to store a moving
87       average of the gradients.
88     """
89     if config is None: config = {}
90     config.setdefault('learning_rate', 1e-2)
91     config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
92     v = config.get('velocity', np.zeros_like(w)) # gets velocity, else sets it to
93     zero.
94     # ===== #
95     # YOUR CODE HERE:
96     # Implement the momentum update formula. Return the updated weights
97     # as next_w, and the updated velocity as v.
98     # ===== #
99     v_new = config['momentum'] * v - config['learning_rate'] * dw
100     next_w = w + v_new + config['momentum'] * (v_new - v)
101     v = v_new
102
103     # ===== #
104     # END YOUR CODE HERE
105     # ===== #
106
107     config['velocity'] = v
108
109     return next_w, config
110
111 def rmsprop(w, dw, config=None):
112     """
113     Uses the RMSProp update rule, which uses a moving average of squared gradient
114     values to set adaptive per-parameter learning rates.
115
116     config format:
117     - learning_rate: Scalar learning rate.
118     - decay_rate: Scalar between 0 and 1 giving the decay rate for the squared
119       gradient cache.
120     - epsilon: Small scalar used for smoothing to avoid dividing by zero.
121     - beta: Moving average of second moments of gradients.
122     """
123     if config is None: config = {}
124     config.setdefault('learning_rate', 1e-2)
125     config.setdefault('decay_rate', 0.99)
126     config.setdefault('epsilon', 1e-8)
127     config.setdefault('a', np.zeros_like(w))
128
129     next_w = None
130
131     # ===== #
132     # YOUR CODE HERE:
133     # Implement RMSProp. Store the next value of w as next_w. You need
134     # to also store in config['a'] the moving average of the second
135     # moment gradients, so they can be used for future gradients. Concretely,
136     # config['a'] corresponds to "a" in the lecture notes.
137     # ===== #
138     config['a'] = config['decay_rate'] * config['a'] + (1 - config['decay_rate']) * dw * dw
139     next_w = w - (config['learning_rate'] * dw) / (np.sqrt(config['a'] + config['epsilon']))
140
141     # ===== #
142     # END YOUR CODE HERE

```

```

143 # ===== #
144
145 return next_w, config
146
147
148 def adam(w, dw, config=None):
149     """
150     Uses the Adam update rule, which incorporates moving averages of both the
151     gradient and its square and a bias correction term.
152
153     config format:
154     - learning_rate: Scalar learning rate.
155     - beta1: Decay rate for moving average of first moment of gradient.
156     - beta2: Decay rate for moving average of second moment of gradient.
157     - epsilon: Small scalar used for smoothing to avoid dividing by zero.
158     - m: Moving average of gradient.
159     - v: Moving average of squared gradient.
160     - t: Iteration number.
161     """
162     if config is None: config = {}
163     config.setdefault('learning_rate', 1e-3)
164     config.setdefault('beta1', 0.9)
165     config.setdefault('beta2', 0.999)
166     config.setdefault('epsilon', 1e-8)
167     config.setdefault('v', np.zeros_like(w))
168     config.setdefault('a', np.zeros_like(w))
169     config.setdefault('t', 0)
170
171     next_w = None
172
173     # ===== #
174     # YOUR CODE HERE:
175     # Implement Adam. Store the next value of w as next_w. You need
176     # to also store in config['a'] the moving average of the second
177     # moment gradients, and in config['v'] the moving average of the
178     # first moments. Finally, store in config['t'] the increasing time.
179     # ===== #
180     config['t'] += 1
181     config['v'] = config['beta1'] * config['v'] + (1 - config['beta1']) * dw
182     config['a'] = config['beta2'] * config['a'] + (1 - config['beta2']) * dw * dw
183     v = config['v'] / (1 - config['beta1'] ** config['t'])
184     a = config['a'] / (1 - config['beta2'] ** config['t'])
185     next_w = w - (config['learning_rate'] * v) / (np.sqrt(a) + config['epsilon'])
186     # ===== #
187     # END YOUR CODE HERE
188     # ===== #
189
190     return next_w, config
191
192
193
194
195
196

```