

# Week 5 Discussion

CS 131 Section 1B

29 April 2022

Danning Yu

# Announcements

- HW3 released, due 5/4
  - Starter code posted under Week 5 on BruinLearn
- Homeworks should be submitted on BruinLearn, under Assignments
- Before submitting
  - Make sure your code compiles on SEASnet server
  - Make sure your function signatures are correct
  - Follow all instructions and specifications
  - Do not submit files in a .zip unless told to do so
- Help and starter code from past TAs
  - <https://github.com/CS131-TA-team>

# Basic Java

Pictures and diagrams borrowed from Boyan Ding and others

# Object Oriented Programming

- Main idea: objects with methods and fields
  - Methods and fields are functions and variables that belong to the object, and encapsulated within
  - Object of the same class share same fields and methods
- Most popular programming paradigm
  - Java, C++, C#, Python, PHP, JavaScript, Ruby, Objective-C, Swift, Scala, Common Lisp, Smalltalk, ...
  - i.e. Most of the popular languages today
- Possible benefits
  - Modularity
  - Information-hiding
  - Code reuse
  - Pluggability and ease of debugging
  - And more

# Classes and Interfaces

- Class: template for an object
  - Object is an instance of a class
  - e.g. We can have multiple Bicycle objects that function the same way, but can be moving at different speed etc.
- All objects created using the same class will have the same methods/fields
- Interface: a description of what needs to be implemented
  - Multiple classes can implement the same interface
  - In Java, a list of functions
  - Allows for separation of API from implementation

# Alan Kay's Definition of OOP

- Everything is an object
  - Numbers, classes, functions, ...
- Objects communicate by sending/receiving messages
  - Think of biological cells communicating
- Objects have their own memory
- Every object is an instance of some class
- All objects of the same type can receive the same messages

**Some of these do not apply to most modern OOP languages!**

# Java

- General-purpose, object-oriented language
- One of the most popular programming languages
- Code compiled into bytecode and runs on a virtual machine
  - What are the pros and cons of this?
- Popular IDEs include Eclipse, IntelliJ IDEA
  - We don't require usage of IDE, you can use any text editor for your homework

# Java: Hello World

- HelloWorld.java

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello, world");  
    }  
}
```

- How to compile

- `javac HelloWorld.java`
- **Generates** `HelloWorld.class` containing bytecode

- Running

- `java HelloWorld`
- **Note:** use class name, not file name



# Files in Java

- `MyClass.java`: **code** for `MyClass`
- `MyClass.class`: **bytecode** for `MyClass` (compiled from `MyClass.java`)
- `Foo.jar`: **Java Archive File**
  - Is really just a ZIP archive
  - Often used to package whole compiled application with resources, configuration, etc
  - Will use this to package source files for Homework 3

# Java Bytecode

- A intermediate form between compiled and interpreted code
  - Platform independence of interpreted code
  - Better performance than interpreted code


```
outer:  
for (int i = 2; i < 1000; i++) {  
    for (int j = 2; j < i; j++) {  
        if (i % j == 0)  
            continue outer;  
    }  
    System.out.println(i);  
}
```

javac

```
0:   iconst_2  
1:   istore_1  
2:   iload_1  
3:   sipush 1000  
6:   if_icmpge 44  
9:   iconst_2  
10:  istore_2  
11:  iload_2  
12:  iload_1  
13:  if_icmpge 31  
16:  iload_1  
17:  iload_2  
18:  irem  
19:  ifne 25  
22:  goto 38  
25:  iinc 2, 1  
28:  goto 11  
31:  getstatic #84;  
34:  iload_1  
35:  invokevirtual #85;  
38:  iinc 1, 1  
41:  goto 2  
44:  return
```

# Java Virtual Machine (JVM)


- Runs bytecode generated by a Java compiler
- Provides separation of code and operating system
  - Allows Java code to run on a variety of OSes
- JVM provides garbage collection, just-in-time compilation (JIT), etc
- Multiple implementations
  - Reference implementation (OpenJDK) provided by Oracle



# OOP In Java

- Abstraction
- Encapsulation
  - Binding data with code that manipulates it
  - Access modifiers: `public/protected/private`
- Inheritance
  - An object may acquire some/all property of another object
- Polymorphism
  - One method can have multiple implementations, usage decided at runtime

# Inheritance in Java



# Inheritance in Java

```
class Shape {  
    void draw() { /* do nothing */ }  
}  
class Rectangle extends Shape {  
    void draw() { /* draw a rectangle */ }  
}  
class Circle extends Shape {  
    void draw() { /* draw a circle */ }  
}  
class Triangle extends Shape {  
    void draw() { /* draw a triangle */ }  
}
```

```
Triangle a = new Triangle();  
/* draws a triangle */  
a.draw();  
  
Shape b = a;  
/* draws a triangle */  
b.draw();  
  
b = new Circle();  
/* draws a circle */  
b.draw();
```

# Inheritance in Java

- Which of the following are allowed?

```
Square a = new Square();  
Shape b = a;
```


```
Shape a = new Shape();  
Square b = a;
```

```
Shape a = new Square();  
Square b = a;
```

- Left: allowed
- Middle: not allowed
  - Shape does not have the same methods or fields as Square
- Right: not allowed
  - Need to cast it

# Inheritance in Java

- Multiple inheritance is not allowed in Java
  - Why?



Diamond Problem



# Interface

- Defines what a class must be able to do, not how to do it
  - Can't be instantiated, should only be implemented by classes
- One class can implement multiple interfaces

```
interface Vehicle {  
    public void increaseSpeed();  
    public void decreaseSpeed();  
    public void turnLeft();  
    public void turnRight();  
}
```

```
class Car implements Vehicle {  
    public void increaseSpeed() {  
        /* Press accelerator */  
    }  
    public void decreaseSpeed() {  
        /* Press brake pedal */  
    }  
    /* other implementations */  
}
```

# Abstract Classes

- Combination of a class and an interface
  - Similar to abstract class in C++ (pure virtual function)
- Objects of an abstract class cannot be created, can only be inherited
- `abstract` method: no implementation, must be implemented by subclasses

```
abstract class Shape {  
    abstract void draw();  
    void setColor(Color c) {  
        /* set color */  
    }  
}
```

# Access Modifiers

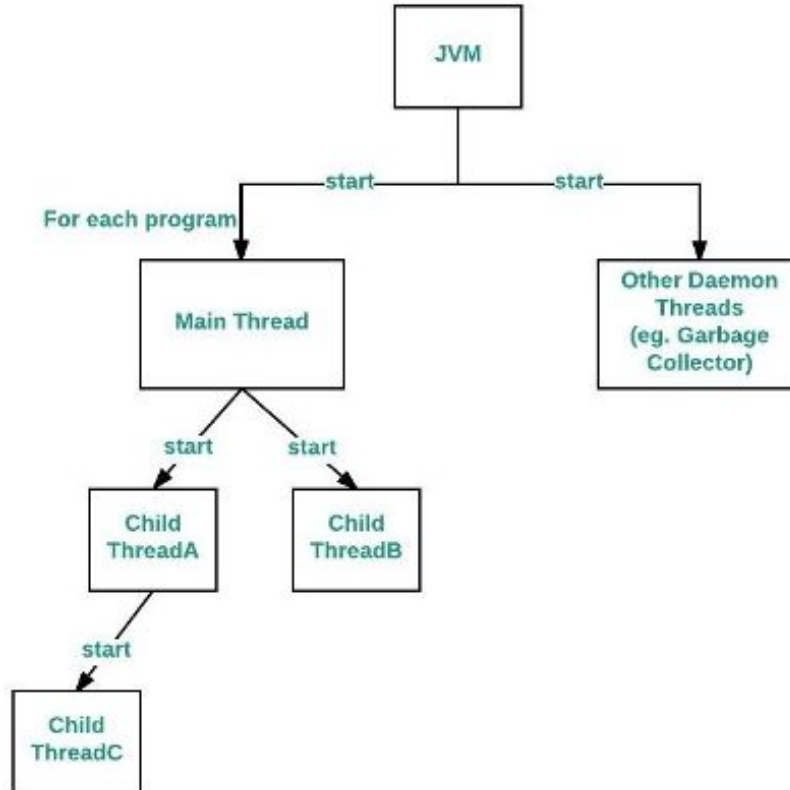
- Controls who can access an object's methods/fields
  - In general, start with `private` and make fields more visible only when necessary
- Classes also have access modifiers: `public` or no modifier (makes it package private)
- Package-private: can access within package but not outside
  - Related Java files are typically grouped together into a package

**Access Levels**

Modifier	Class	Package	Subclass	World
<code>public</code>	Y	Y	Y	Y
<code>protected</code>	Y	Y	Y	N
<i>no modifier</i>	Y	Y	N	N
<code>private</code>	Y	N	N	N

# Java Memory Model

# Threads in Java



# Creating and Using Threads

```
public class MyRunnable implements Runnable {  
    public void run() {  
        System.out.println("MyRunnable - START ");  
        // Do some heavy processing here  
        System.out.println("MyRunnable - END ");  
    }  
}
```

// In your main method:

```
Thread t1 = new Thread(new MyRunnable());  
Thread t2 = new Thread(new MyRunnable());  
t1.start(); // Start executing thread 1  
t2.start(); // Start executing thread 2  
t1.join(); // Wait for thread 1 to finish  
t2.join(); // Wait for thread 2 to finish
```

# Java Memory Model

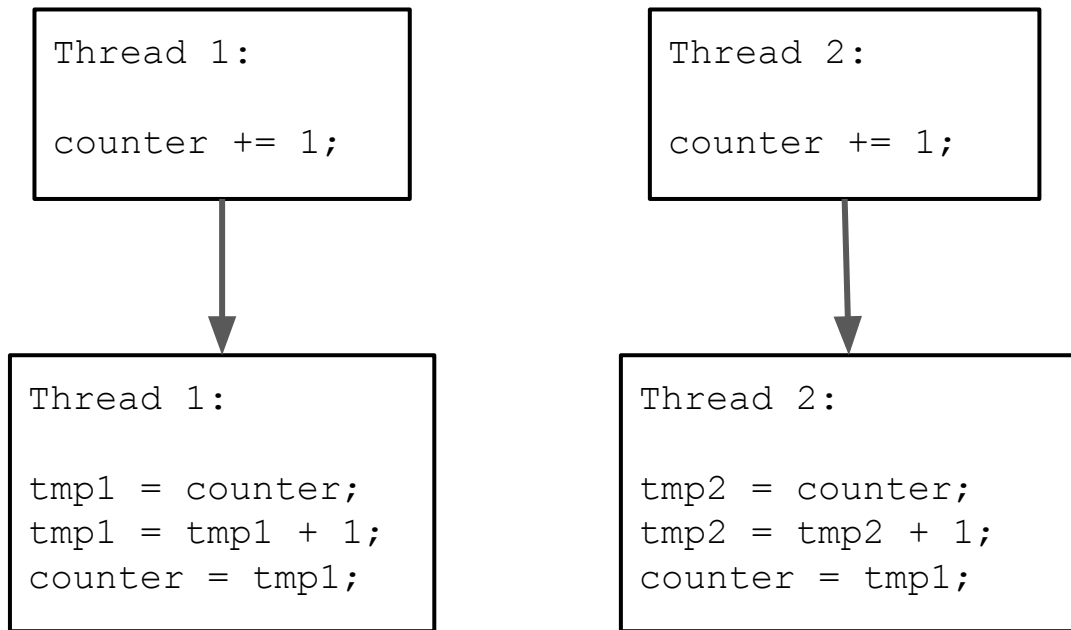
- Defines how threads interact with memory
- Defines what code reorderings are legal for the compiler and processor to carry out
- “As-if-serial” semantics used within one thread
  - Compiler can change your code in any way as long as the result of execution is the same
- e.g. in the following code, `x` and `y` are both initially 0
  - Java may freely reorder the assignment on the left, but not on the right

```
x = 2;  
y = 1;  
System.out.format("%d %d\n", x, y);
```

```
x = 2;  
System.out.format("%d %d\n", x, y);  
y = 1;  
System.out.format("%d %d\n", x, y);
```

- With multiple threads are running, the situation gets complicated

# Race Conditions





# synchronized Keyword

- Each object has a lock
- This keyword enforces exclusive access
  - Only one thread can enter a `synchronized` method in one object at once
- Happens-before relationship
  - Everything that one thread did while in a synchronized block will be visible to the next thread entering a synchronized block
- A thread can call any other synchronized methods while it holds the lock

# synchronized **Keyword**: Example

```
public class SynchronizedCounter {  
    private int c = 0;  
  
    public synchronized void increment() {  
        c++;  
    }  
  
    public synchronized void decrement() {  
        c--;  
    }  
  
    public synchronized int value() {  
        return c;  
    }  
}
```

## synchronized Keyword: Example

- Synchronized can also be used with smaller blocks of code
- Avoid blocking others when it's not necessary

```
public class SynchronizedCounter {  
    private int c = 0;  
  
    public void incrementAndWork() {  
        // ... computation here ....  
        synchronized(this) {  
            c++;  
        }  
        // ... computation here ....  
    }  
}
```

# synchronized Keyword

- Any object can be used as the lock for synchronized

```
public class MyClass {  
    private int c1 = 0;  
    private int c2 = 0;  
    private Object lock1 = new Object();  
    private Object lock2 = new Object();  
    public void inc1() {  
        synchronized(lock1) {  
            c1++;  
        }  
    }  
    public void inc2() {  
        synchronized(lock2) {  
            c2++;  
        }  
    }  
}
```

# volatile Keyword

- Guarantees that other threads will see the changes immediately
  - `volatile` access can not be reordered relative to other reads/writes
  - Effectively serves as a memory barrier
- Excellent explanation with details [found at this link](#)

Thread 1:

```
x = 5;  
done = true;
```

Thread 2:

```
while (!done) {}  
System.out.println(x);
```

- Without `volatile`, the printed value may not always be 5
- If `done` is defined as `volatile`, then `x` will always be printed as 5

# Atomic Operations

- Another option for preventing race conditions
- Atomic operations: code is translated into assembly instructions that guarantee the update value is visible to all threads without race conditions
  - Typically means a single assembly instruction, but sometimes more
  - Other threads do not see intermediate states, only final state
- Atomic package `java.util.concurrent.atomic` provides data types with atomic operations
- `AtomicInteger` can be used to perform `cnt++` as an atomic operation

```
AtomicInteger cnt = new AtomicInteger(5);  
cnt.incrementAndGet();
```

# Optional: C++ Memory Model

- C++ did not have a memory model until C++11
  - Before then, multithreading behavior was technically unspecified!
- Great talk on the C++ memory model
  - Part 1: <https://www.youtube.com/watch?v=A8eCGOqgvH4>
  - Part 2: <https://www.youtube.com/watch?v=KeLBd2EJLOU>

# Homework 3



# Introduction

- Multithreaded gzip compression
  - Implementing pigz in Java
  - No need to implement the actual compression algorithm, but rather, make the existing compression algorithm multithreaded
- Can leverage the ideas of an existing implementation called [MessAdmin](#)
- Also use starter code on BruinLearn to help you
- Test and compare 4 programs: gzip and pigz in Linux; Pigzj on JVM and native version (compiled with native-image)

# GZip Compression

- gzip: a stream compression format based on DEFLATE
- Input and output are binary streams
- Commonly used for file compression
  - .gz extension
  - Note: .tar is just a tarball: a collection of files (with no compression) usage: e.g. file compression
- Format specified in RFC 1952
- A gzip file can have one or more members shown on the right
- Java implements this algorithm and other compression algorithms in `java.util.zip`

Each member has the following structure:

```
+---+---+---+---+---+---+---+---+
| ID1|ID2|CM |FLG|      MTIME      |XFL|OS | (more-->)
+---+---+---+---+---+---+---+---
```

Header

(if FLG.FEXTRA set)

```
+---+---+---+---+---+---+---+---+
| XLEN |...XLEN bytes of "extra field"...| (more-->)
+---+---+---+---+---+---+---+---
```

(if FLG.FNAME set)

```
+---+---+---+---+---+---+---+---+
|...original file name, zero-terminated...| (more-->)
+---+---+---+---+---+---+---+---
```

(if FLG.FCOMMENT set)

```
+---+---+---+---+---+---+---+---+
|...file comment, zero-terminated...| (more-->)
+---+---+---+---+---+---+---+---
```

Header extensions

(if FLG.FHCRC set)

```
+---+---+
| CRC16 |
+---+---+
```

```
+---+---+---+---+---+---+---+---+
|...compressed blocks...| (more-->)
+---+---+---+---+---+---+---+---
```

Compressed Data

```
      0   1   2   3   4   5   6   7
+---+---+---+---+---+---+---+---+
|      CRC32      |      ISIZE      |
+---+---+---+---+---+---+---+---
```

Tail (checksum + size)

# GZip Compression in Java

- Implementing a simple gzip equivalent in Java
- Take input from stdin, writes output to stdout
  - Equivalent to `gzip -c -`

```
import java.io.IOException;
import java.util.zip.GZIPOutputStream;

public class SimpleGZip {
    public static void main(String[] args) throws IOException {
        GZIPOutputStream gzout = new GZIPOutputStream(System.out);
        System.in.transferTo(gzout);
        gzout.close();
    }
}
```

# Parallel GZip Compression

- A basic implementation of gzip process the file linearly
- How to parallelize compression (with  $P$  threads)
  - Break the files into  $P$  partitions, with each thread processing one partition. Then, concatenated the compressed partitions together
  - Pigz's approach: divide input into fixed size blocks (128 KiB), and have  $P$  threads busily processing a block
    - 1 KiB = 1024 bytes
- What's the difference? Why do we prefer the latter for the HW?
  - Allows us to handle streaming cases
  - As each thread finishes, append to the end of the file (fast)

# Pigz Details

- From pigz's manual page:
  - Checksum: "The individual check value for each chunk is also calculated in parallel... A combined check value is calculated from the individual check values." (note: not implemented in MessAdmin)
  - Dictionary: "The input blocks, while compressed independently, have the last 32K of the previous block loaded as a preset dictionary to preserve the compression effectiveness..."

# Java Environment

- Make sure you're using the correct version of Java as specified in the HW specifications
- Can also try installing GraalVM on your own machine
  - <https://www.graalvm.org>
  - You need the GraalVM and Native Image tool
- The native-image tool from GraalVM allows people to compile Java programs to native binaries
  - To use it, first compile Java code to class files as usual
  - Then use “native-image ClassName” to compile the java program, and you will get an executable named “classname”
  - It can be quite slow on SEASnet, use it when you are sure your java program is functioning correctly with JVM

# MessAdmin

- Github link: <https://github.com/MessAdmin/MessAdmin-Core>
  - Relevant code can be found in  
`clime.messadmin.utils.compress.{gzip,impl}` packages
- The final implementation will need a class similar to [PGZIPOutputStream](#)
- You can learn the ideas and techniques used there, especially the following
  - The [Compressor](#) class, which uses `ThreadPoolExecutor` for multithreaded compression execution
  - The style of passing tasks among threads
- You're welcome to write code in your own style if you want

# Homework Requirements

- The main program should be named `Pigzj`
  - This means your main class is called `Pigzj`
  - Optionally takes argument `-p processes` to specify the threads used, default to number of processors on the system
  - Input taken from `stdin`, output goes to `stdout`, no need for file operation
- Other requirements
  - Correctness: your compressed file should be understood by `gzip/pigz`
  - For full credit, output should only contain a single member
  - Ideally, the output should be byte-for-byte identical with `pigz` output
  - Give proper error messages for certain cases
- Full requirements found on homework page



# Homework Submission

- Submit a single jar file containing
  - All the `.java` source files
  - Do not add compiled `.class` files
- A report with performance measurements and analysis
  - Plain text file under 60 kB
  - Compare the runtime of `gzip`, `pigz`, and `Pigzj` (both JVM and GraalVM native) with different settings
  - Try to use `strace` and see if the result can explain your findings
  - Check HW specification for full list of requirements
- Make sure to use the provided commands to sanity check your submission

Thank You