

# CS 161 Intro. To Artificial Intelligence

Week 2, Discussion 1B

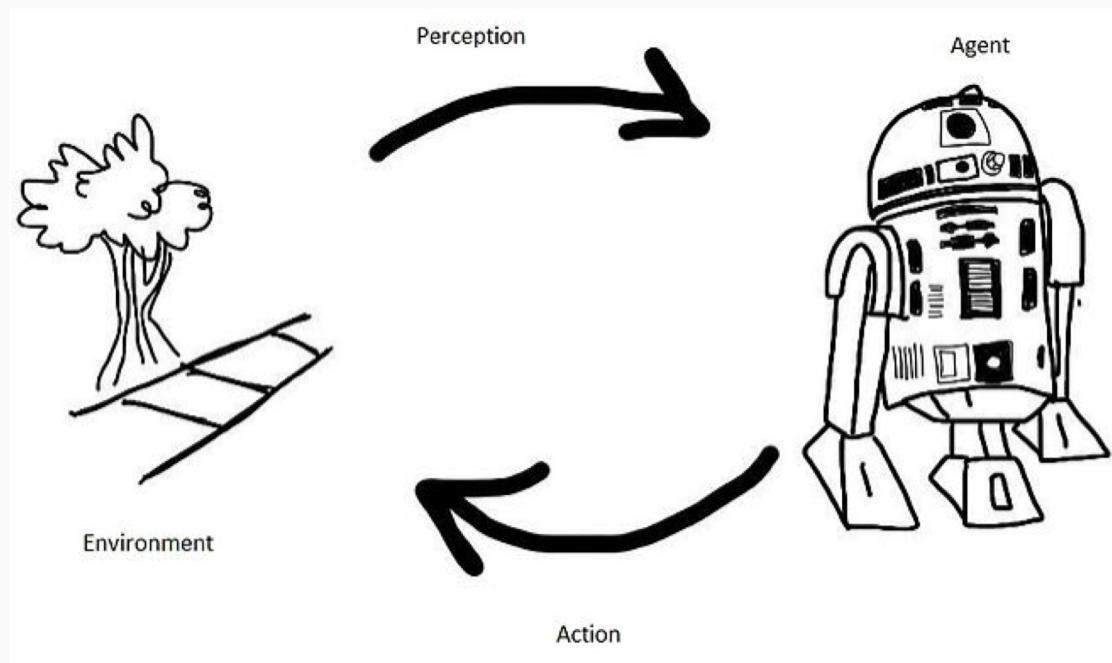
Qian Long

# Today's Topics:

- Agents
- Search Problem Formulation
- State Space Graph & Search Tree
- Search Algorithm Evaluation
- Uninformed Search
  - Breadth-first Search (BFS)
  - Uniform-cost Search (UCS)
  - Depth-first Search (DFS)
  - Depth-limited Search
  - Iterative Deepening Search

# Agents

In AI, an intelligent agent **perceives** its **environment** through **sensors** and **acts** upon it through **actuators**.



# Agents

2 types of agents:

- **Rational agents:**
  - Choose actions based on maximized expected utility
    - Goal-based agents (eg. reach a goal with lowest cost)
    - Utility-based agents (eg. reinforcement learning)
- **Reflex agents:**
  - Choose actions based on current percept of the environment
  - Does not consider future consequence of actions
  - If-else condition-action: if current condition  $\Rightarrow$  action
    - Eg. a mail sorting robot (see an address  $\Rightarrow$  put letter into a right bag)

# Search Problem Formulation

A search problem consists of:

- **Initial state**
- **State space:**  $S = \{s_1, s_2, \dots, s_d\}$
- **Actions (operators):** a set of possible actions
- **Transition model (successor function):**  $F(s_t, a_t) = s_{t+1}$ , sometimes with a path cost function
- **Goal test:** determine if the solution/goal is achieved

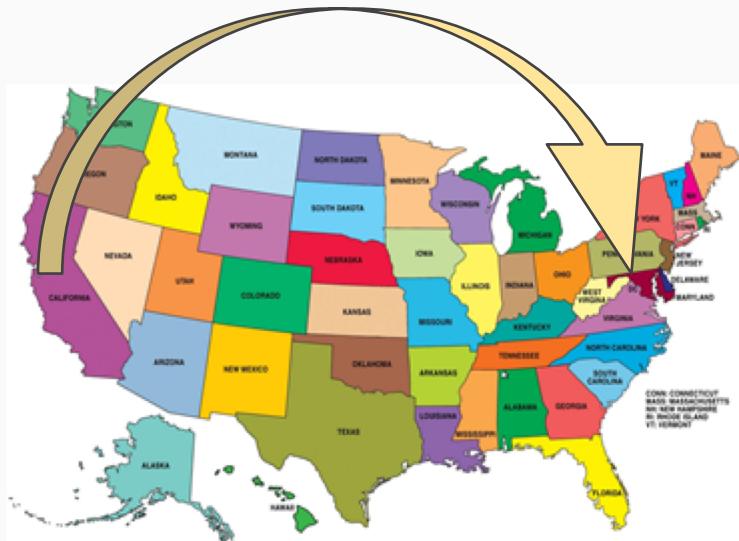
A **solution:** a sequence of actions that transform the initial state to a goal state

**Problem formulation:** the process of deciding what actions and states to consider, given a goal.

# Search Problem Formulation - Example

## Travel from CA to MA:

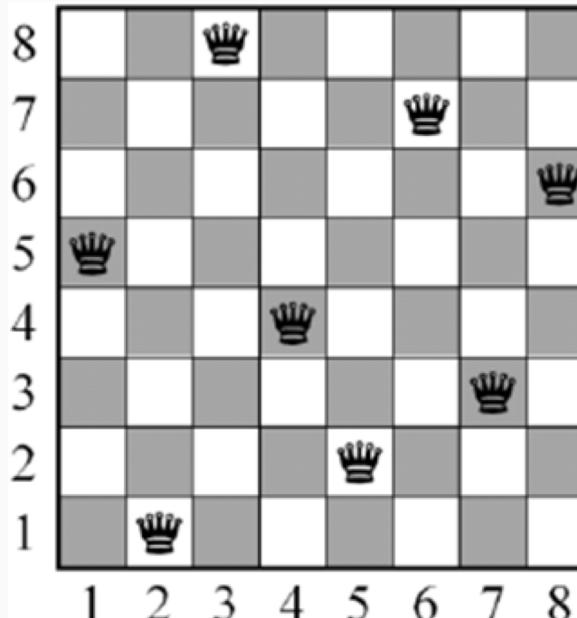
- Initial State: CA
  - State space: 50 states in U.S.
  - Actions: go to adjacent state (cost=distance)
    - Eg. Actions(CA) = {GoTo(OR), GoTo(AZ), GoTo(NV)}
  - Successor function (transition model):
    - Eg. Result(Ini(CA), GoTo(AZ)) = Ini(AZ)
  - Path cost function:
    - Eg. sum of dist, # of actions, none, etc.
  - Goal test: state == MA?



# Search Problem Formulation - Example

The 8 Queens puzzle:

- **Objective:** Place eight queens on a chessboard such that no queen attacks any other.  
⇒ No two queens on the same row, column, diagonal



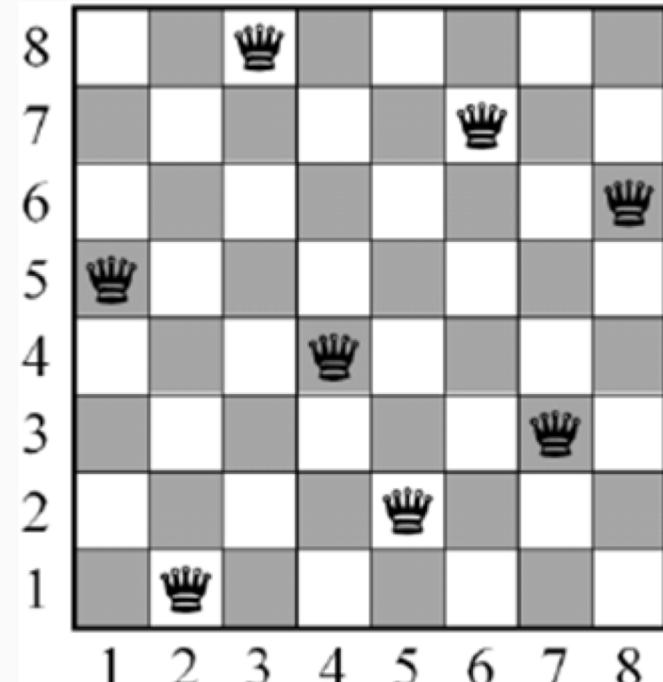
# Search Problem Formulation - Example

The 8 Queens puzzle (version 1):

- Initial State: no queen on the board
- State space: any arrangement of 0 to 8 queens on the board is a state
- Actions: add a queen to any empty square
- Successor function (transition model): returns the board with a queen added to the specified square
- Goal test: 8 queens on the board, none attack each other

This is an **incremental formulation**.

Q: How large is state space?

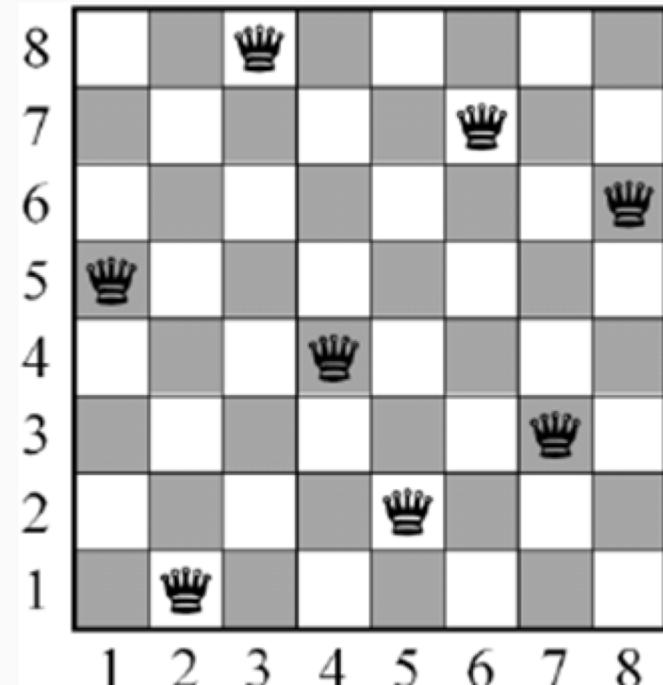


# Search Problem Formulation - Example

The 8 Queens puzzle (version 2):

- Initial State: no queen on the board
- State space: all possible arrangements of  $n$  queens ( $0 \leq n \leq 8$ ), one per column in the leftmost  $n$  columns, with no queen attacking another
- Actions: add a queen to any square in the leftmost empty column such that it is not attacked by any other queen
- Successor function (transition model): returns the board with a queen added to the specified square
- Goal test: 8 queens on the board, none attack each other

State space reduced from  $1.8 * 10^{14}$  to 2057 (given in textbook)!



# Search Problem Formulation - Example

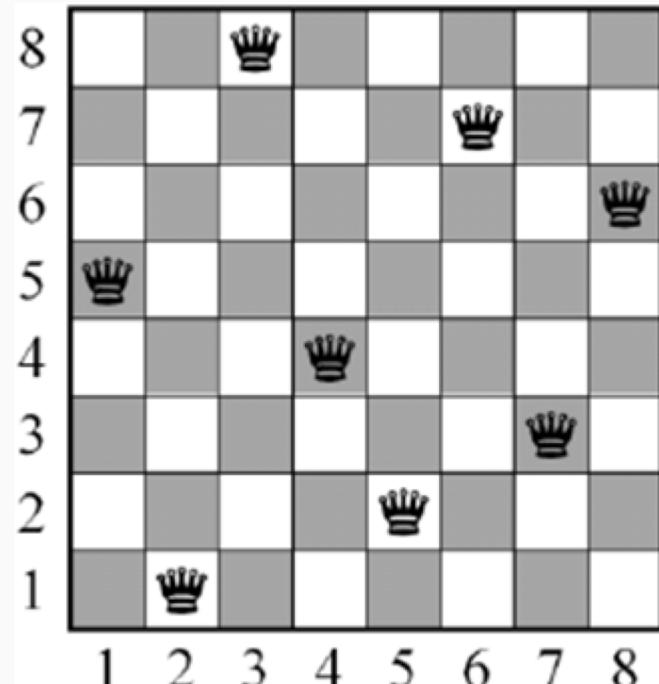
## Incremental formulation

- Start with an empty state
- Each action adds a queen

## Complete-state formulation

- Start with all 8 queens on the board
- Moves each queens around

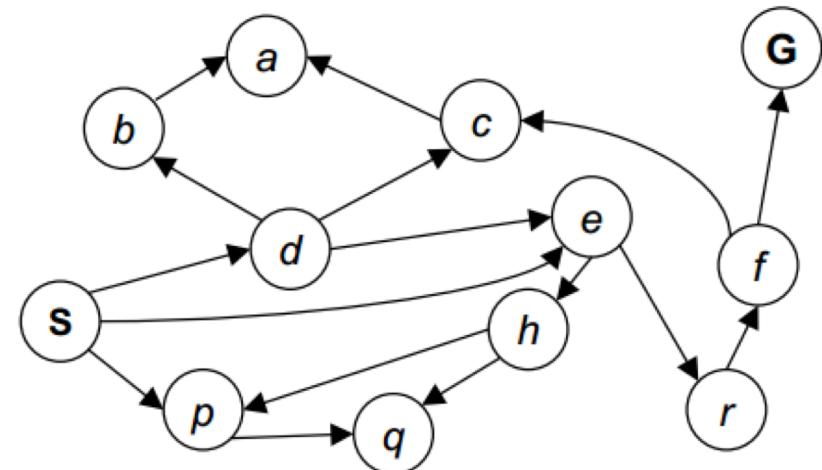
Note: In either case, path cost doesn't matter because in our problem only final state counts.



# State Space Graph

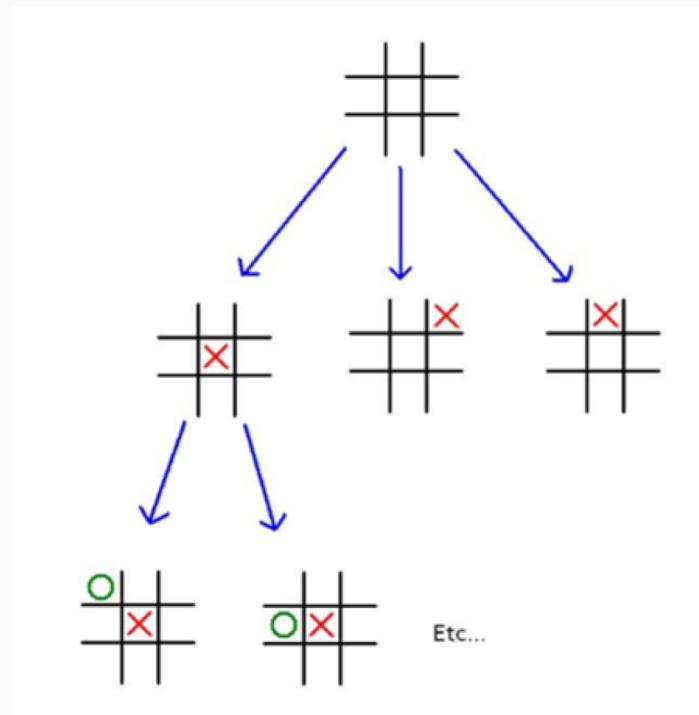
**State space graph** is a mathematical representation of a search problem

- Nodes are (abstracted) world configurations, represent states
- Arcs represent successors (action results)
- Goal test: one or a set of goal nodes
- Each state occurs only once!



# Search Tree - Basics

- A "what if" tree of plans and their outcomes
- Root is initial state
- Children correspond to successor states  
(correspond to plans to those states)



# Search Tree Algorithm

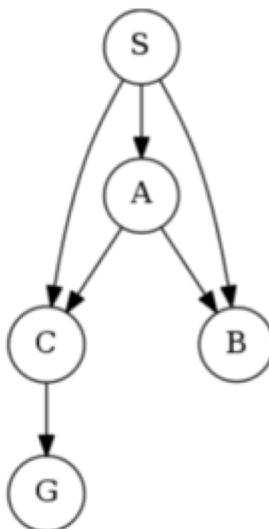
- **Basic idea:** offline, simulated exploration of state space by generating successors of already-explored states
- Strategy: Breadth-first Search (BFS), Depth-first Search (DFS), etc. (will go through these later)

```
function TREE-SEARCH( problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

Goal test

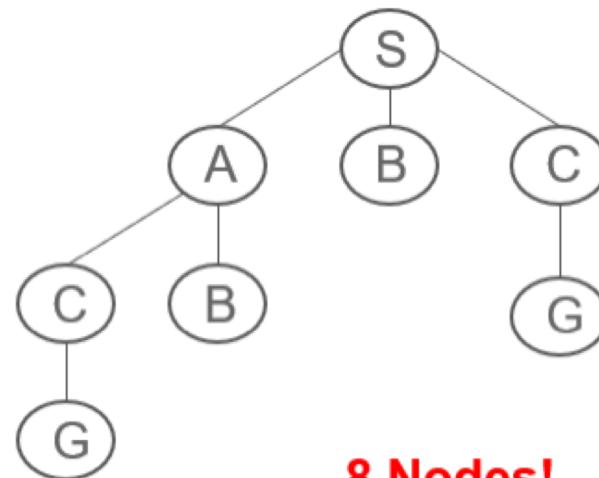
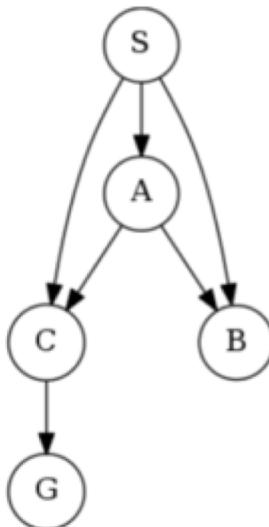
# State Space Graph & Search Tree - Example

Given the following state graph, start from S and end with G, how many nodes are there in this search tree?



# State Space Graph & Search Tree - Example

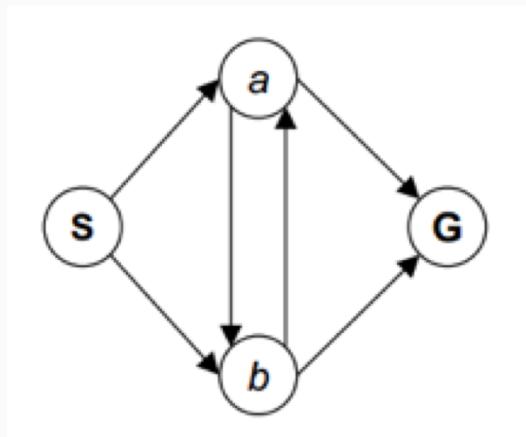
Given the following state graph, start from S and end with G, how many nodes are there in this search tree?



**8 Nodes!  
(including the root node)**

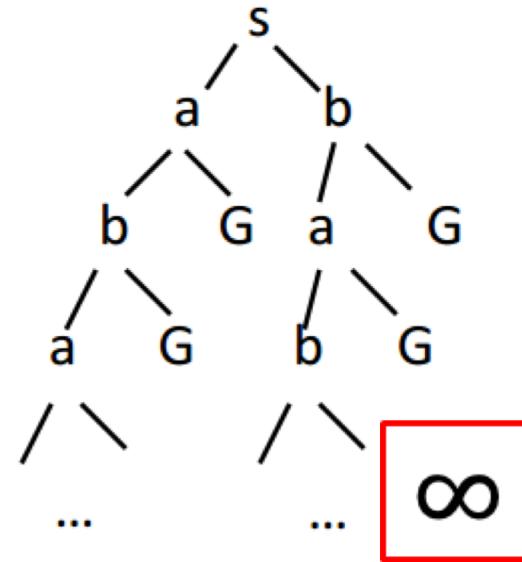
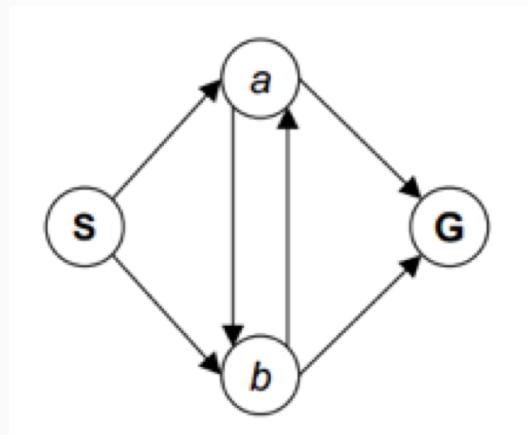
# State Space Graph & Search Tree - Example

Given the following 4-state graph, start from S and end with G, how many nodes are there in this search tree?



# State Space Graph & Search Tree - Example

Given the following 4-state graph, start from S and end with G, how many nodes are there in this search tree?



# Search Algorithm Evaluation

- **Completeness**
  - Does it always find a solution if one exists?
- **Optimality**
  - Does it always find a least-cost solution?
    - In search trees, this is often in terms of depth of the solution if no path cost given
    - For some algorithm (eg. Uniform-cost Search), it can be in terms of path cost
- **Time Complexity**
  - Number of nodes generated/expanded
- **Space Complexity**
  - Maximum number of nodes in memory

# Search Algorithm Evaluation

Time and space complexity are measured in terms of:

- **b**: maximum # of children nodes for one parent node, also called **branching factor**
- **d**: depth of the least cost solution
- **m**: maximum depth of the state space (may be infinity)

Other useful terms:

- **Frontier/Fringe**: the (generated) nodes to be expanded
- **Generation**: create a node and have it in memory
- **Expansion**: pop a node from frontier and generate its children (expand)

# Uninformed (Blind) Search

- Breadth-first Search (BFS)
- Uniform-cost Search (UCS)
- Depth-first Search (DFS)
- Depth-limited Search
- Iterative Deepening Search

# Breadth-first Search (BFS)

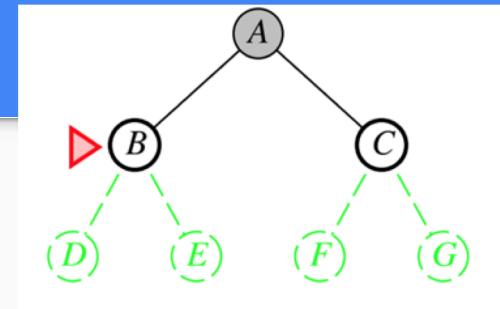
Expands **shallowest** nodes first

We can do **goal test with generation or expansion**:

- If goal test with **generation**:
  - root node A is give, it is not count into one layer
  - when expand A, we generate B and C
  - **then we goal test B and C**
  - then we expand B and generate D and E
  - **then we goal test D and E**
- If goal test with **expansion**:
  - root node A is give, it is not count into one layer
  - when expand A, we generate B and C
  - **then we goal test A**
  - then we expand B and generate D and E
  - **then we goal test B**

Suppose our goal is C:

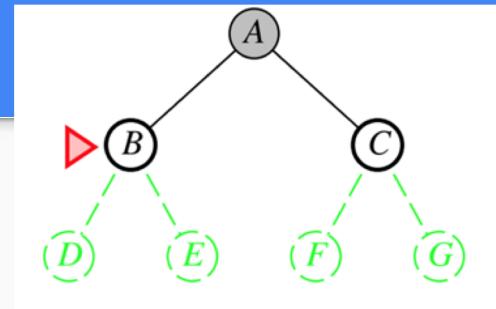
- When goal test with generation, after we generate C, we stop
- When goal test with expansion, after we generate C, it's not enough, we need to continue after we generate F and G
- **By default, we do goal test with expansion for BFS**



# Breadth-first Search (BFS)

Evaluation:

- Complete (if  $b$  is finite)
- Optimal for unit step costs (i.e. each step's cost is same)
- Time complexity:
  - If goal test during the generation:  $b+b^2+\dots+b^d=O(b^d) \rightarrow$  by default
  - If we goal test on expansion:  $b+b^2+b^3+\dots+b^{d+1}=O(b^{d+1})$
- Space complexity is  $O(b^d)$ 
  - we have to keep all nodes in memory
  - $1+b+\dots+b^d$  (including root node)



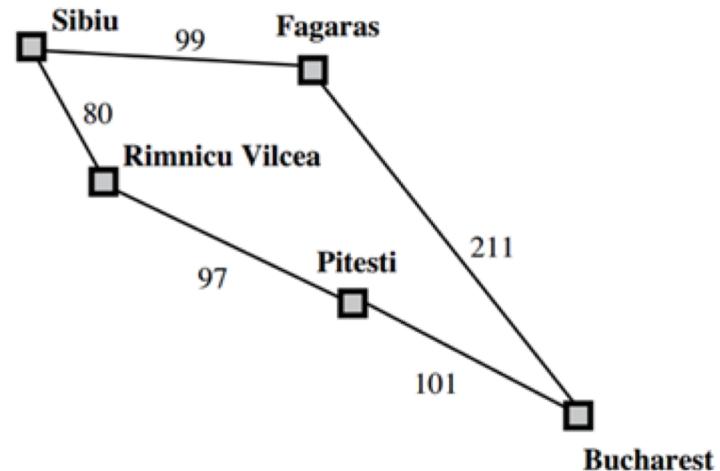
# Uniform-cost Search (UCS)

Expands the node with **lowest path cost**

UCS keeps going after goal node has been generated, and **terminate after a goal node is expanded**.

Example: From Sibiu (S) to Bucharest (B)

- Start with S, expand it and get RV and F
- **Goal test S**
- Then, since RV cost 80, F cost 99, so we expand RV and get P
- **Goal test RV**
- Then, since F cost 99, P cost  $80+97=177$ , so we expand F and get B (**and we continue!!**)
- **Goal test F**
- Then, since B cost  $99+211=310$ , P cost  $80+97=177$ , so we expand P and get B again
- **Goal test P**
- Then we expand and do goal test for B. It has 2 path, we choose the path with minimum cost



# Uniform-cost Search (UCS)

Evaluation:

- Complete (if  $b$  is finite and the step costs has a positive lower bound  $\epsilon$ )
- Optimal in general (**cost = cumulative path cost instead of depth for UCS**)
- Time complexity:
  - $C$ : cost (NOT depth) of optimal solution
  - Every action costs at least  $\epsilon$
  - Worst case time complexity:  $O(b^{1+\lfloor C/\epsilon \rfloor})$ 
    - If all step costs are equal:  $O(b^{1+d})$  (similar to BFS but test on expansion)
- Space complexity:
  - Fringe: priority queue (where priority is cumulative cost)
  - Worst case space complexity: roughly the last tier,  $O(b^{1+\lfloor C/\epsilon \rfloor})$

# Uniform-cost Search & BFS

## BFS:

- **Stops** after a goal node is expanded (unless otherwise specified)
- Optimal when we have unit step cost

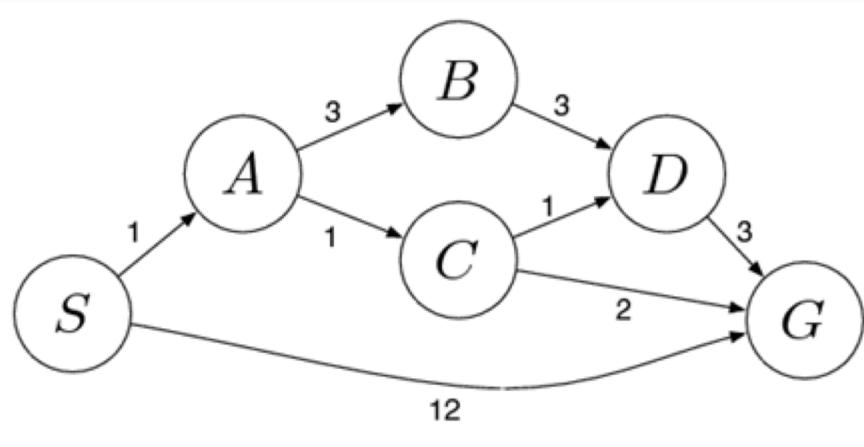
## UCS:

- **Keeps going** after a goal node has been generated (since it do goal test when expansion)
- Can deal with the case that all step costs are not equal
- When we have unit step cost, UCS is similar to BFS (except the timing of goal test)

# Uniform-cost Search (UCS) - Example

Use uniform-cost search (assume S is initial state, G is goal state)

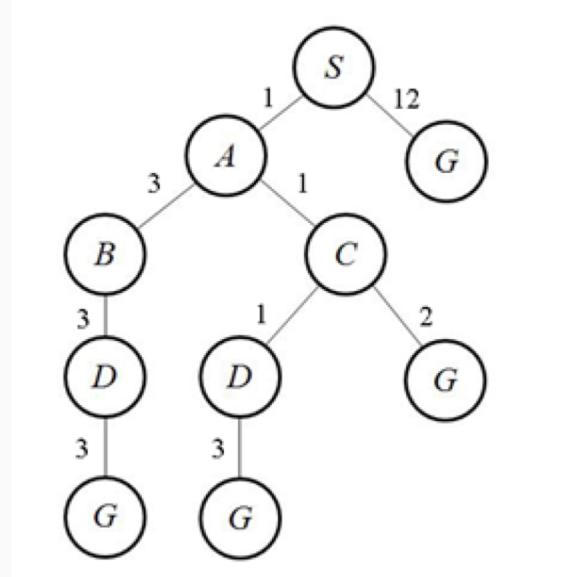
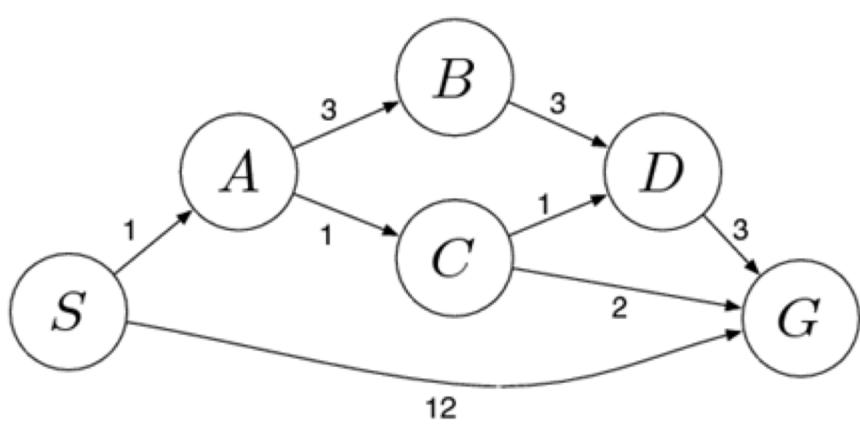
- Give the generated search tree?
- Show in what order we expand nodes (choose left-to-right when have tie)?
- Return the optimal solution/path?



# Uniform-cost Search (UCS) - Example

Use uniform-cost search (assume S is initial state, G is goal state)

- Give the generated search tree?
- Show in what order we expand nodes (choose left-to-right when have tie)?  
A: S->A->C->D->B->G
- Return the optimal solution/path?  
A: S->A->C->G, cost is 4

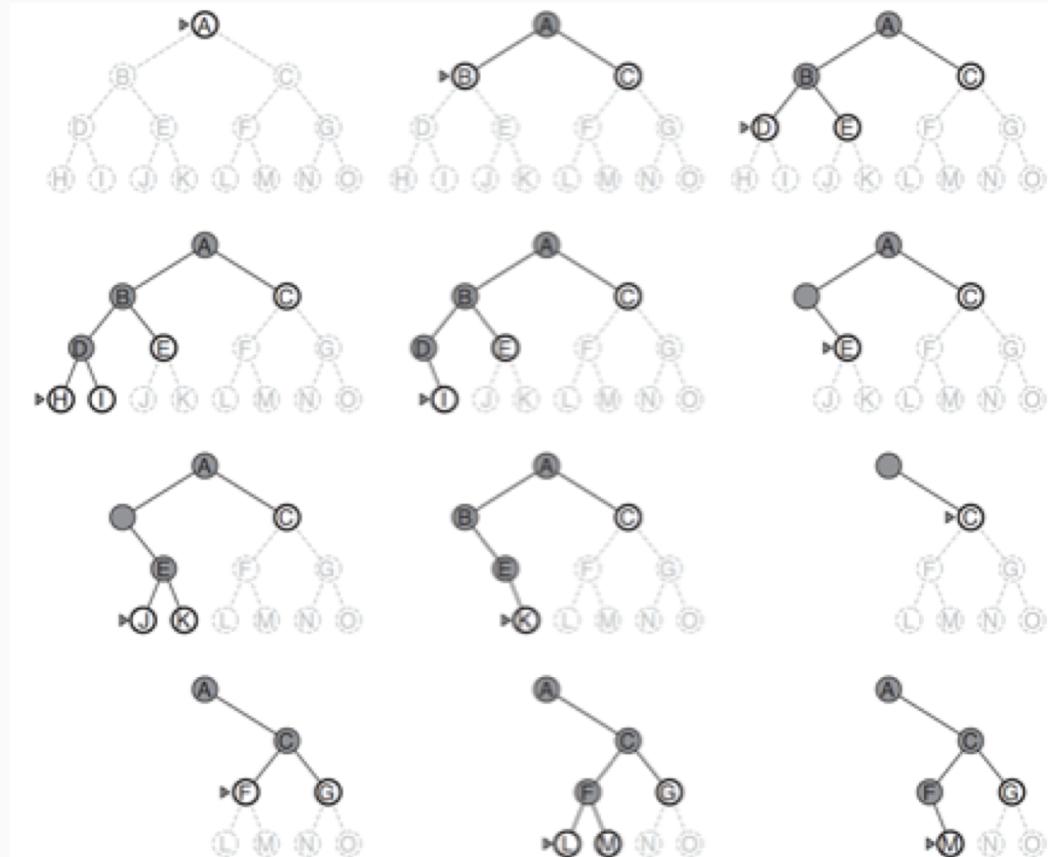


# Depth-first Search (DFS)

Expand **deepest** unexpanded node

Example:

- M is target, shallowest solution  
depth  $d = 3$
- State space max depth  $m = 3$
- Light grey nodes: unexplored region
- Grey nodes: expanded nodes
- White nodes: generated but not expanded node  $\Rightarrow$  frontier / fringe
- Removed nodes: explored nodes with no descendants in the frontier



# Depth-first Search (DFS)

Evaluation:

- Not complete (a tree can be unbounded)
- Not optimal (if  $m > d$ )
- Time complexity:  $O(b^m)$ , where  $m$  is maximum depth
  - $1+b+b^2+\dots+b^m$
  - If tree is unbounded,  $m$  is infinite
- Space complexity:  $O(bm) \rightarrow$  linear space!
  - Fringe = nodes on current search path + siblings along the path
  - # nodes keep in memory:  $b+b+\dots+b$  ( $m$  times)

# Depth-limited Search

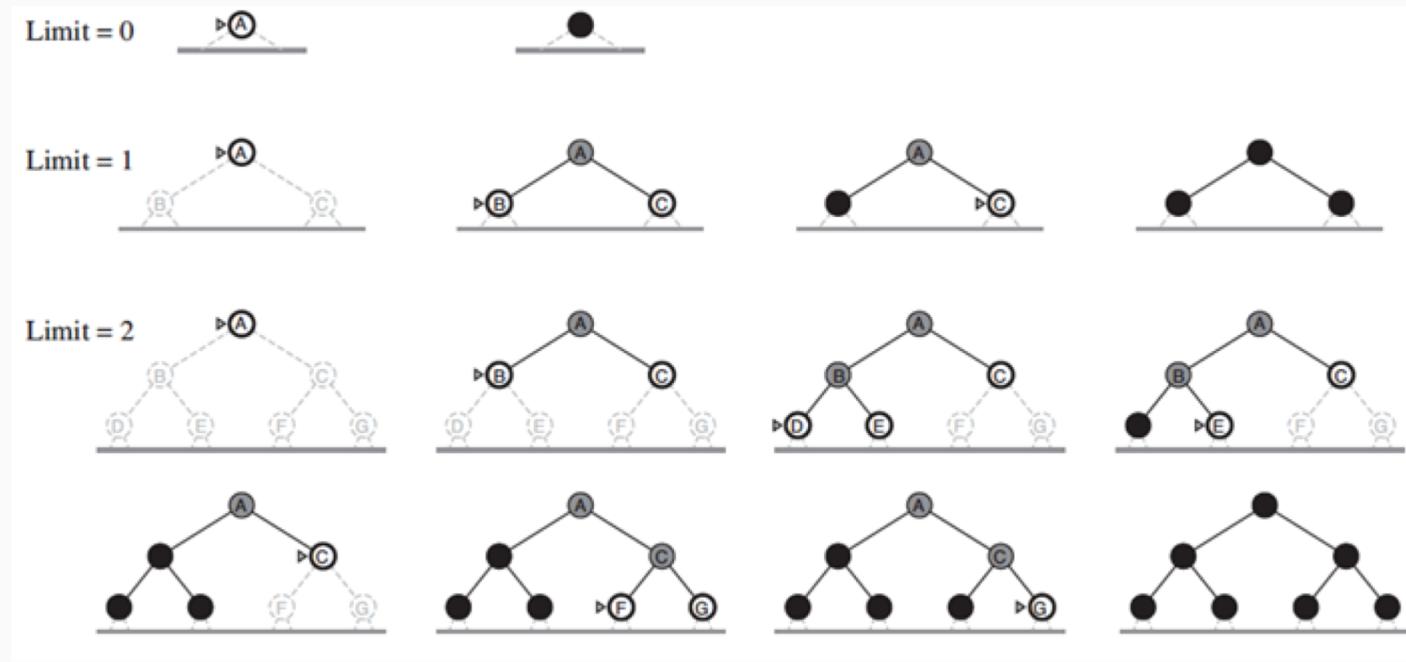
- Compare to DFS, it add a depth limit  $L$  (often written in lowercase  $l$ )
- Even if the tree can be further expanded, we stopped when reach depth  $L$

Evaluation:

- Not complete in general, only complete if  $L \geq d$
- Not optimal (same as DFS)
- Time complexity:  $O(b^L)$ , where  $L$  is the depth limit
- Space complexity is  $O(bL)$

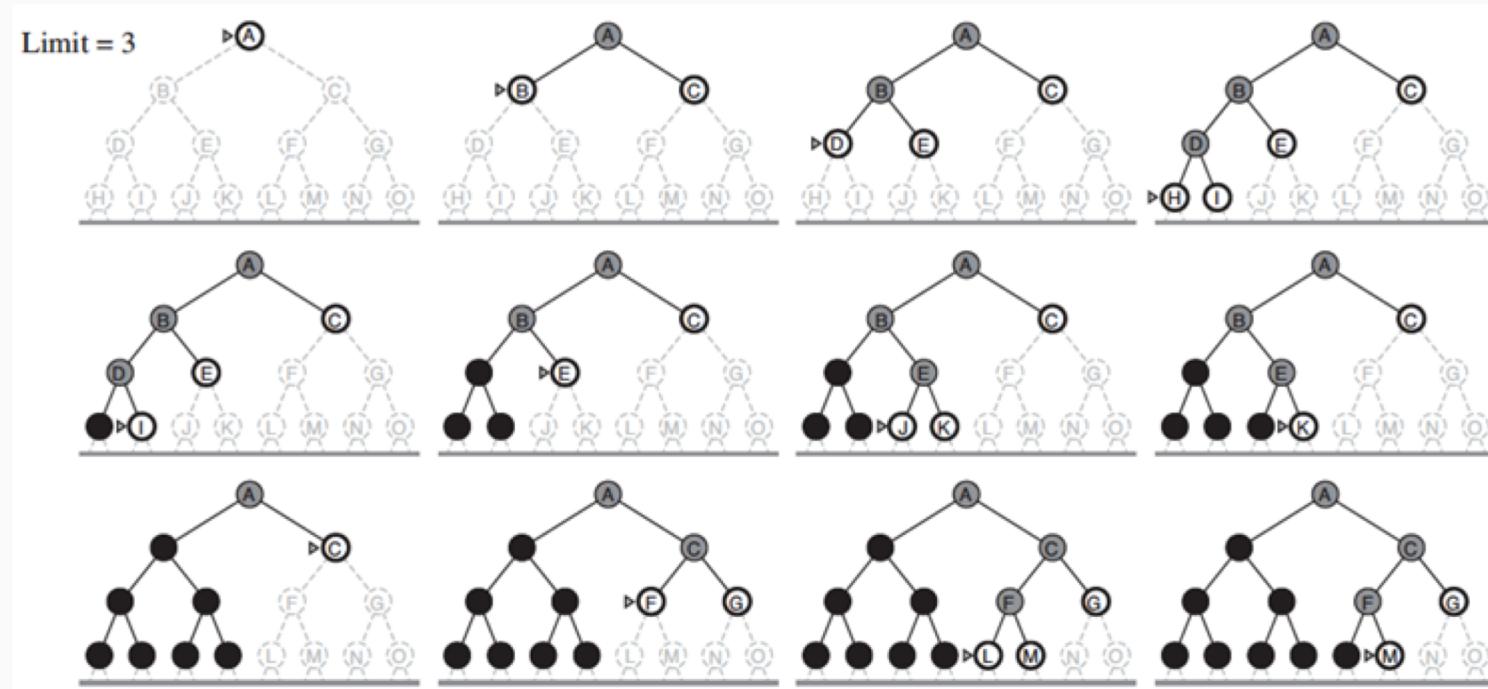
# Iterative Deepening Search (ID / IDS)

- Similar to DFS and Depth-limited Search given a particular depth limit
- **Increase depth limits** until a goal is found



# Iterative Deepening Search (ID / IDS)

- Similar to DFS and Depth-limited Search, but **increase depth limits** until a goal is found



# Iterative Deepening Search (ID / IDS)

Evaluation:

- Complete (if  $b$  is finite)
- Optimal for unit step costs (i.e. each step's cost is same)
- Time complexity:  $O(b^d)$ , where  $d$  is the depth of the shallowest solution
  - $db + (d-1)b^2 + \dots + b^d$ , where the last item  $b^d$  dominates
  - This is because the children of root have been generated  $d$  times (when depth limit  $l$  goes from 1 to  $d$ , we need to generate them), then the children of those children of root have been generated  $d-1$  times, and so on...
- Space complexity is  $O(bd)$

# Summary of Uninformed Search

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes <sup>a</sup>	Yes <sup>a,b</sup>	No	No	Yes <sup>a</sup>
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$
Optimal?	Yes <sup>c</sup>	Yes	No	No	Yes <sup>c</sup>

**Figure 3.21** Evaluation of tree-search strategies.  $b$  is the branching factor;  $d$  is the depth of the shallowest solution;  $m$  is the maximum depth of the search tree;  $\ell$  is the depth limit. Superscript caveats are as follows: <sup>a</sup> complete if  $b$  is finite; <sup>b</sup> complete if step costs  $\geq \epsilon$  for positive  $\epsilon$ ; <sup>c</sup> optimal if step costs are all identical; <sup>d</sup> if both directions use breadth-first search.

# Questions?

- My slides take the following materials as references:
  - Xiao Zeng's slides

Thank you!